

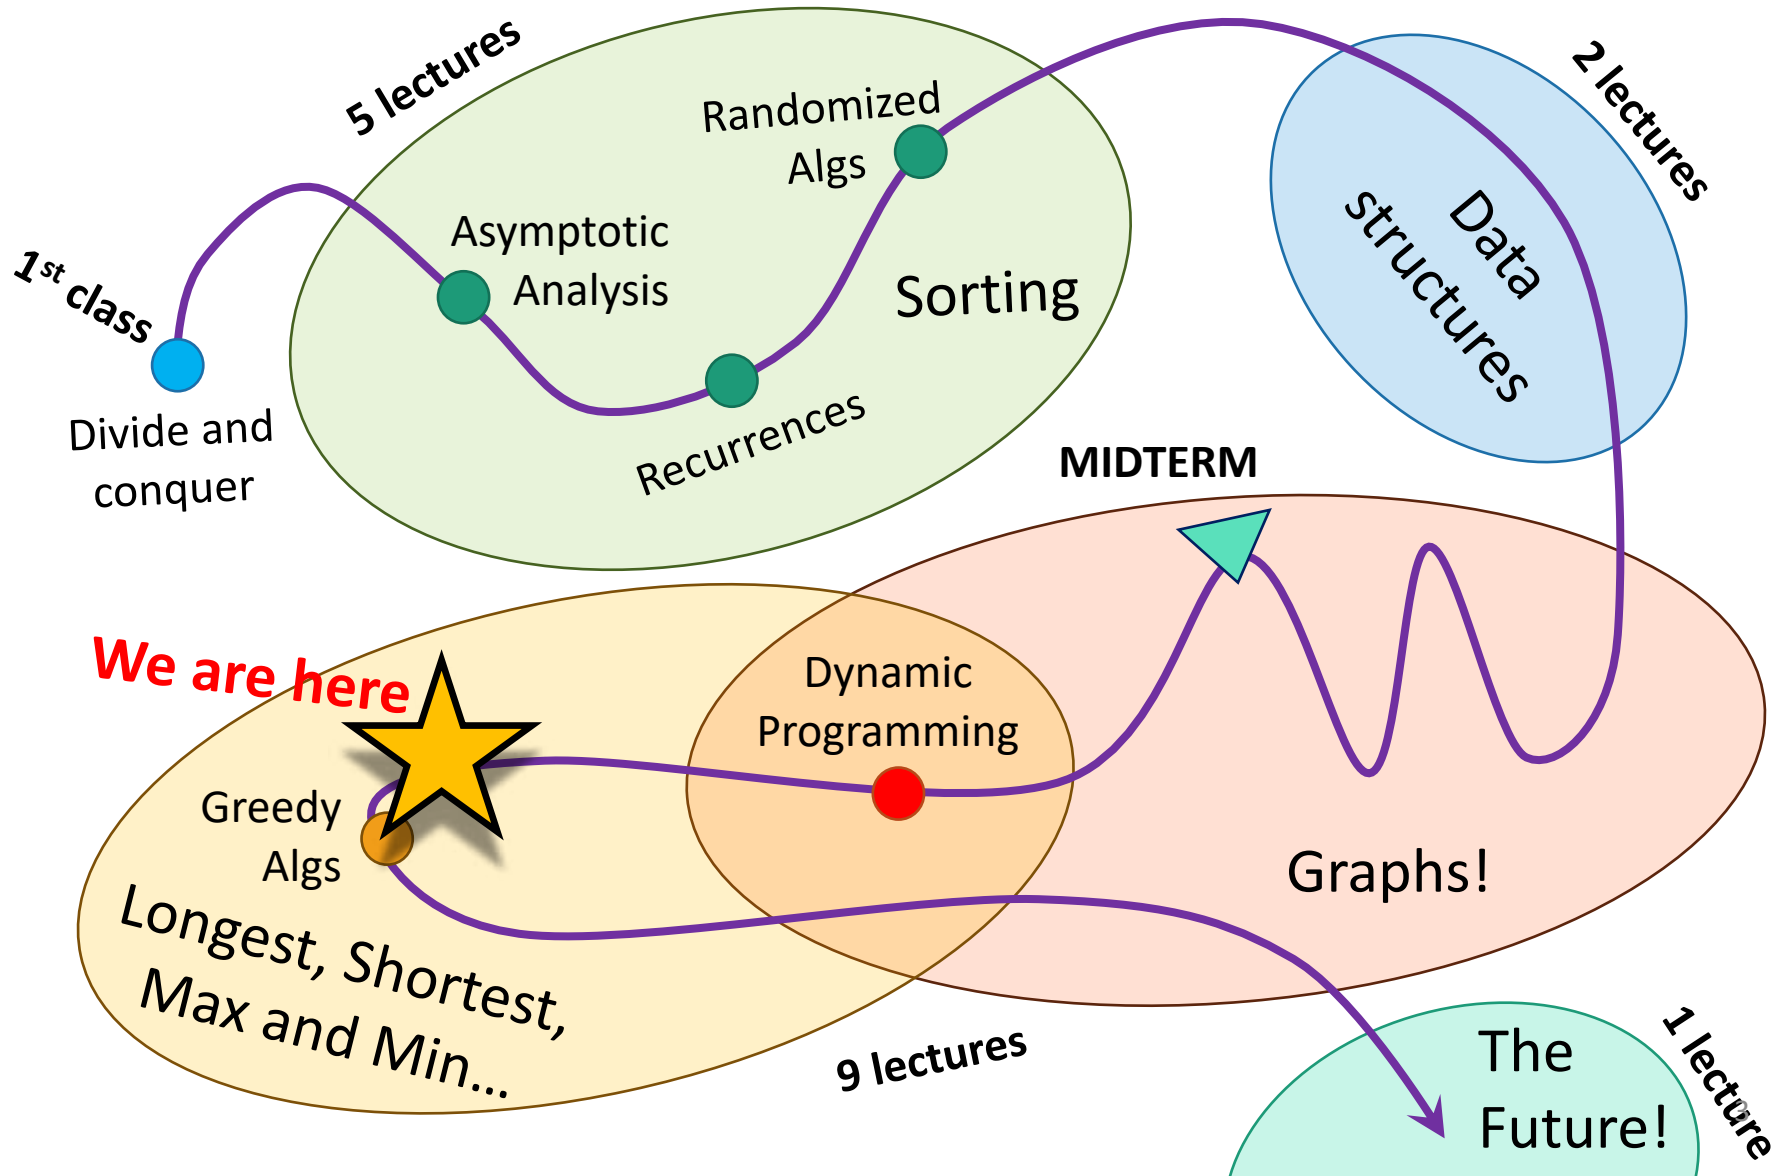
Lecture 14

Greedy algorithms!

Announcements

- HW6 Due Friday!

Roadmap



This week

- Greedy algorithms!

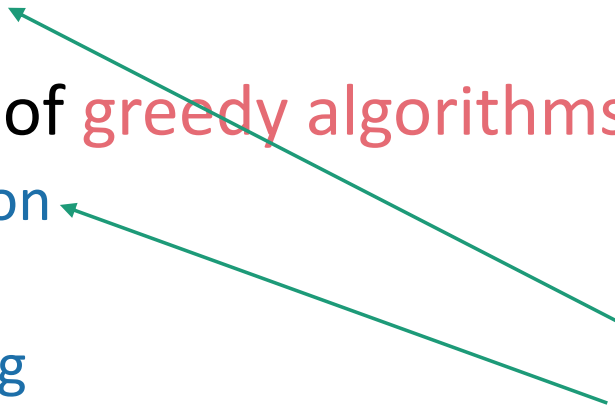


Greedy algorithms

- Make choices one-at-a-time.
- Never look back.
- Hope for the best.

Today

- One example of a **greedy algorithm** that **does not work**:
 - Knapsack again
- Three examples of **greedy algorithms** that **do work**:
 - Activity Selection
 - Job Scheduling
 - Huffman Coding



You saw these on
your pre-lecture
exercise!

Non-example

- Unbounded Knapsack.



Capacity: 10

Item:



Weight:

6

2

4

3

11

Value:

20

8

14

13

35

- Unbounded Knapsack:

- Suppose I have **infinite copies** of all of the items.
- What's the **most valuable way to fill the knapsack?**



Total weight: 10

Total value: 42

- **“Greedy”** algorithm for unbounded knapsack:

- Tacos have the best Value/Weight ratio!
- Keep grabbing tacos!



Total weight: 9

Total value: 39

Example where greedy works

Activity selection

You can only do one activity at a time, and you want to maximize the number of activities that you do.

What to choose?

CS 161 Class

Math 51 Class

Sleep

CS 161
Section

nar

Progra
team

5 Class

Underwater basket
weaving class

CS110
Class

CS161 study
group

Swimming
lessons

Theory Lunch

Social activity


Combinatorics
Seminar

time

Activity selection

- Input:


- Activities a_1, a_2, \dots, a_n
- Start times s_1, s_2, \dots, s_n
- Finish times f_1, f_2, \dots, f_n



- Output:

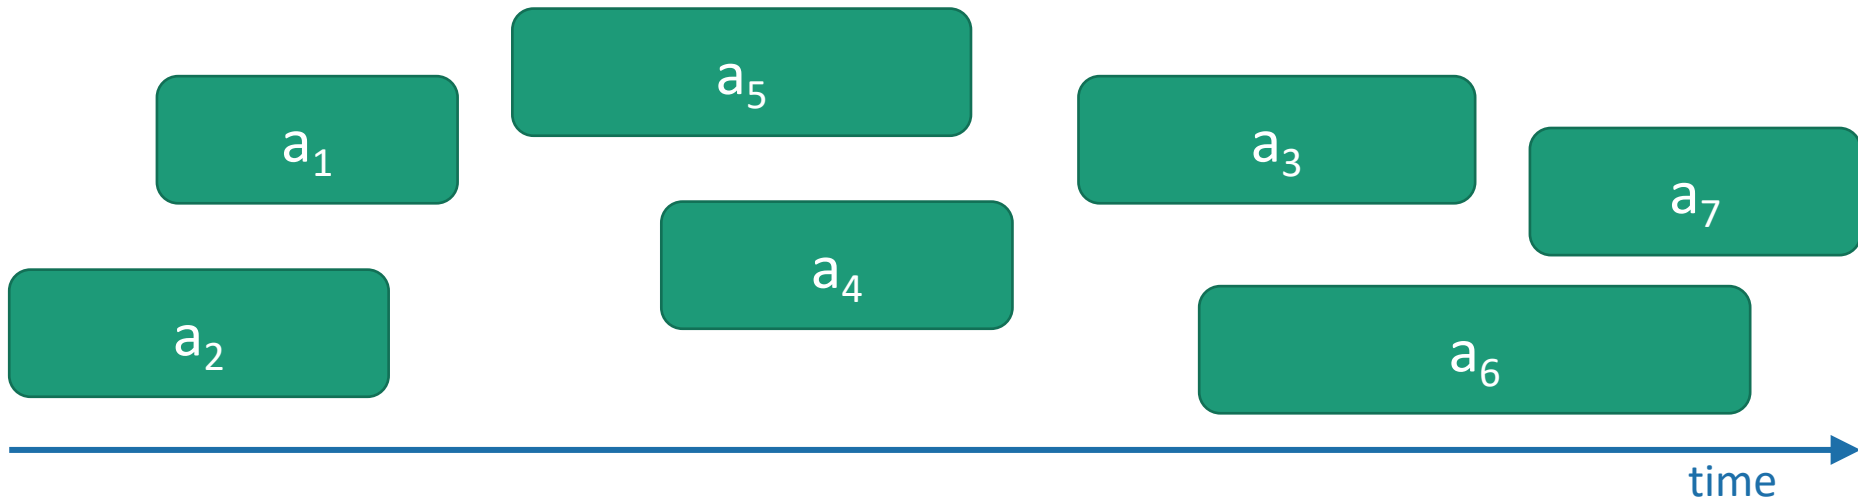
- A way to maximize the number of activities you can do today.

In what order should you greedily add activities?



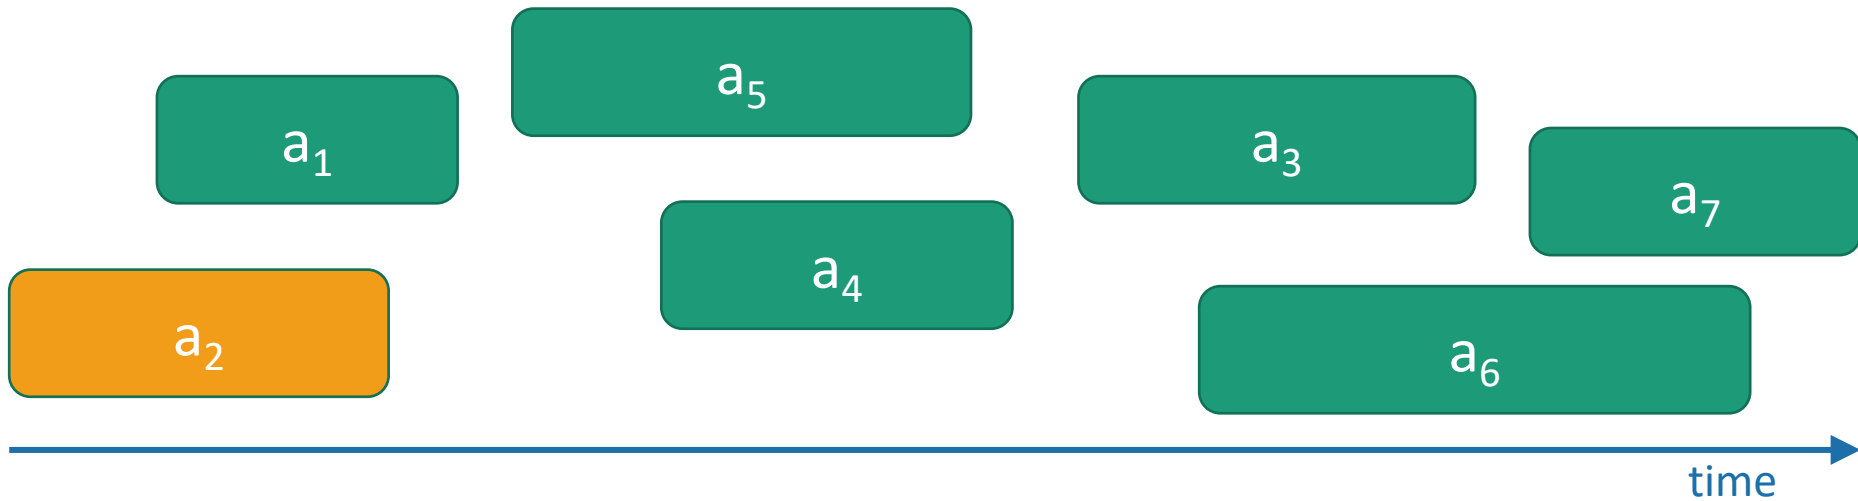
Think-pair-share!

Greedy Algorithm



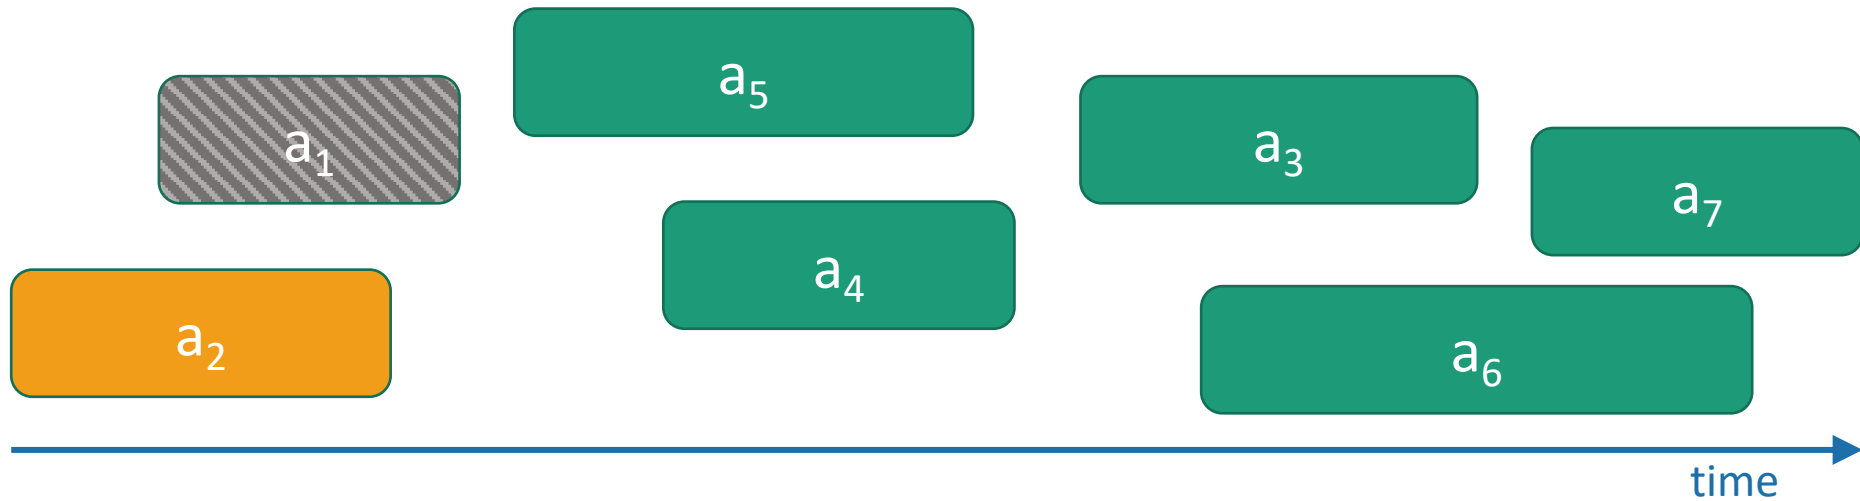
- Pick activity you can add with the smallest finish time.
- Repeat.

Greedy Algorithm



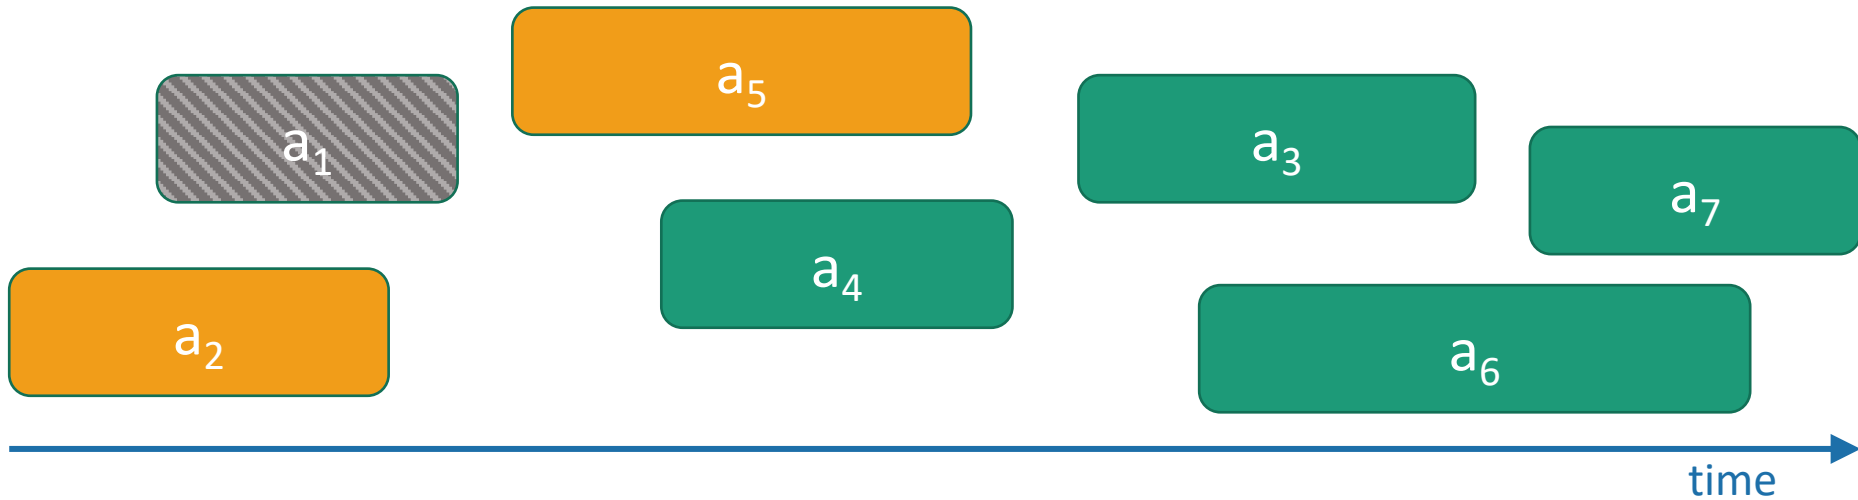
- Pick activity you can add with the smallest finish time.
- Repeat.

Greedy Algorithm



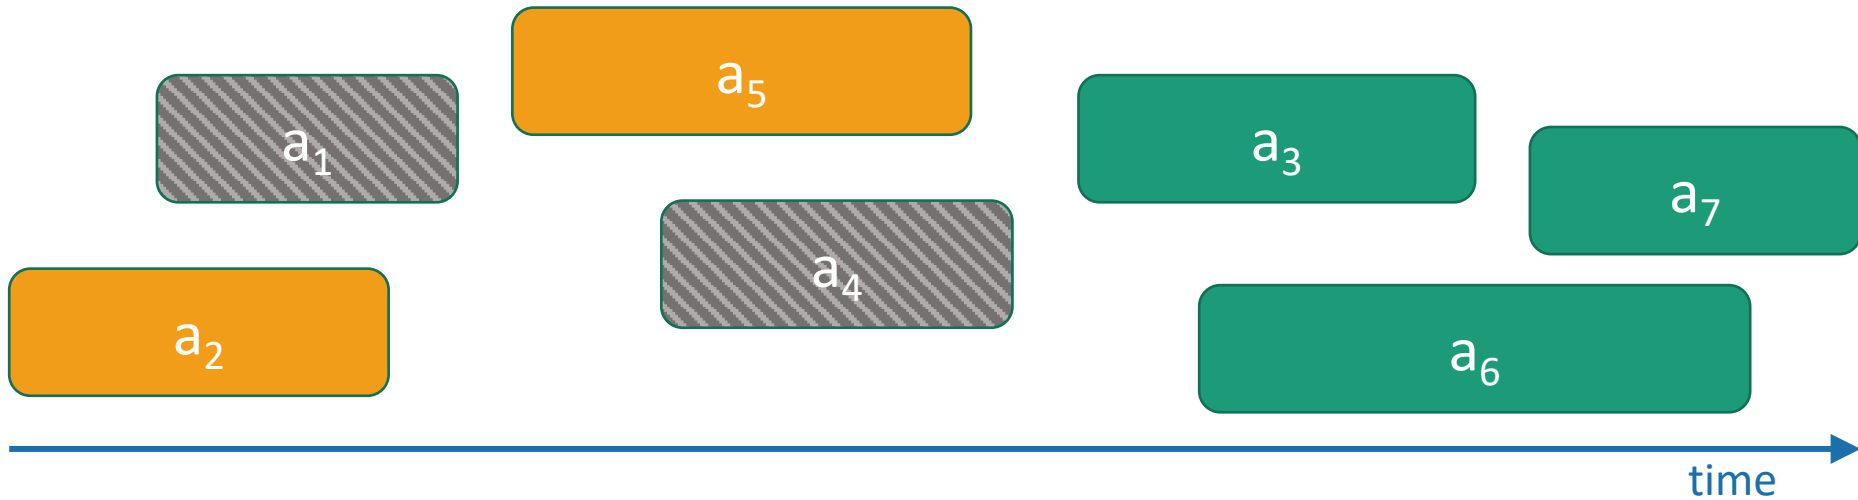
- Pick activity you can add with the smallest finish time.
- Repeat.

Greedy Algorithm



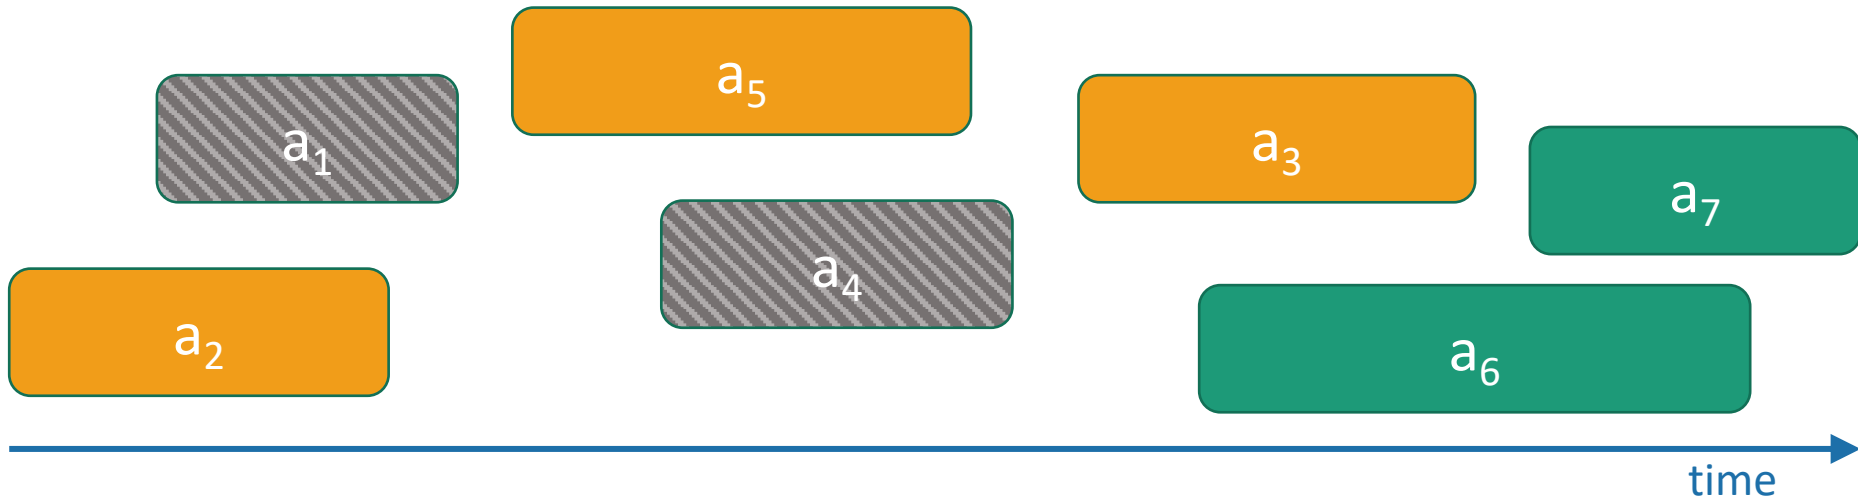
- Pick activity you can add with the smallest finish time.
- Repeat.

Greedy Algorithm



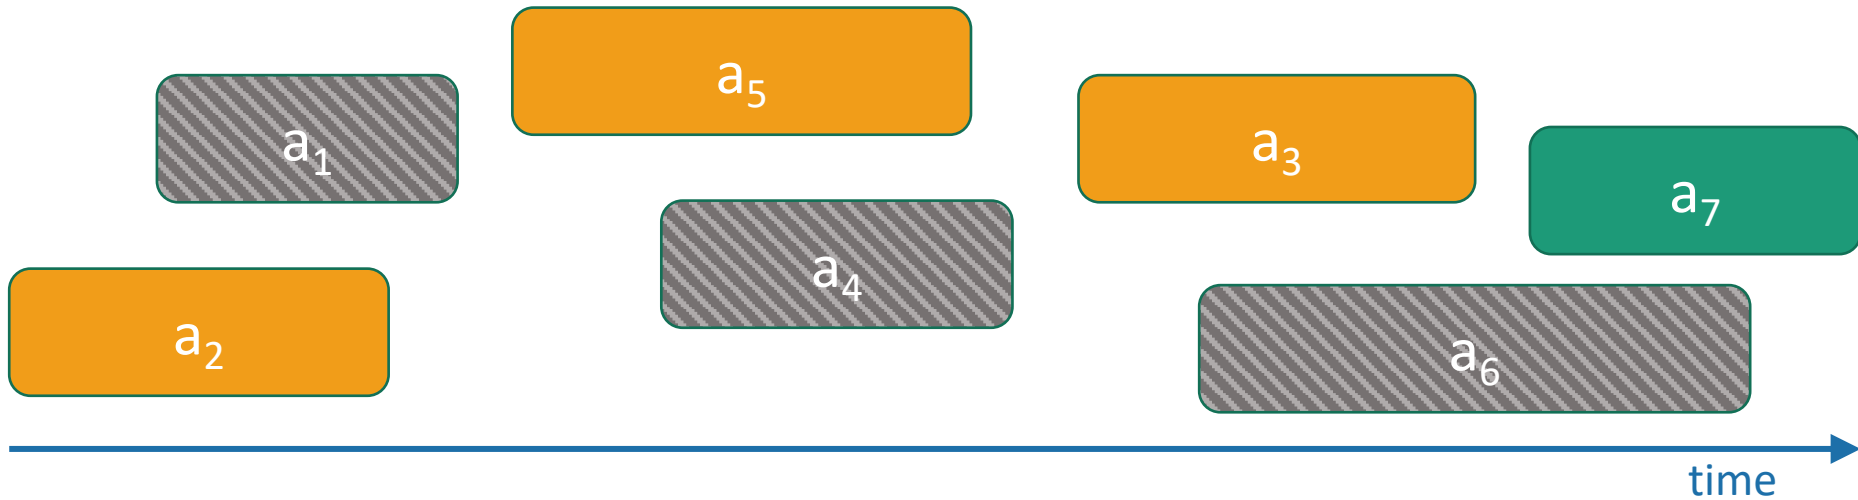
- Pick activity you can add with the smallest finish time.
- Repeat.

Greedy Algorithm



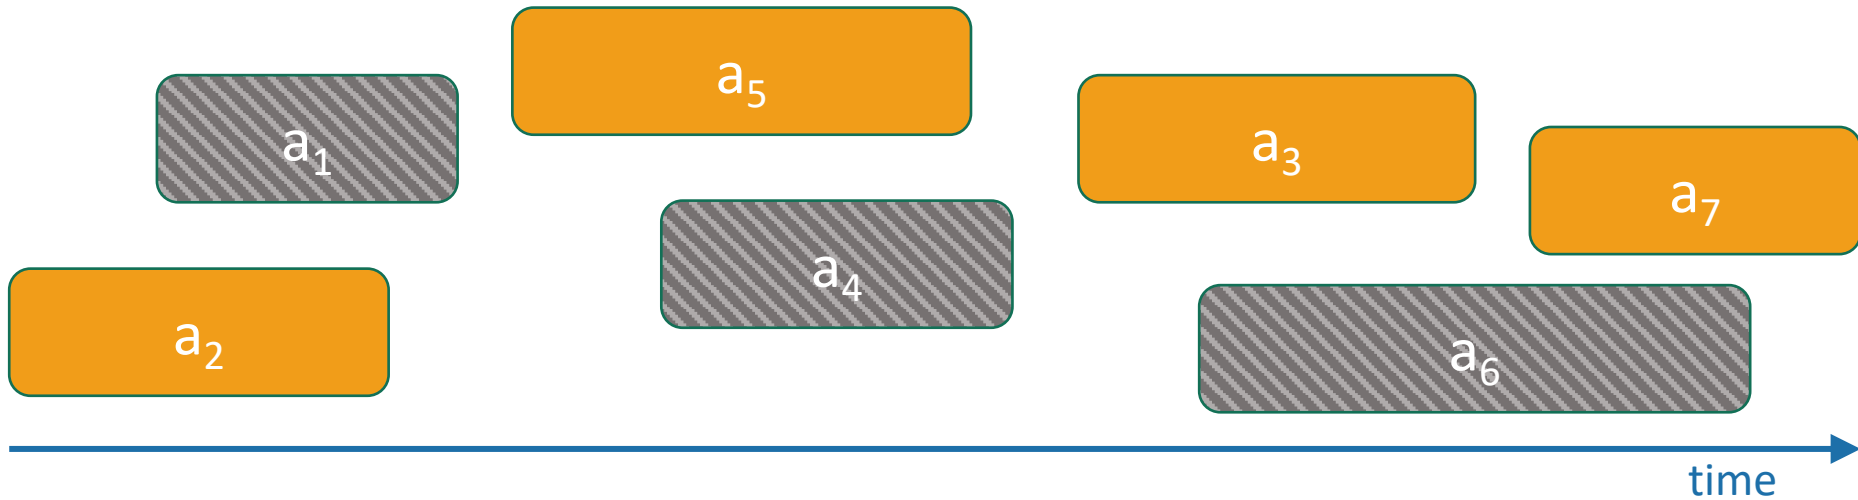
- Pick activity you can add with the smallest finish time.
- Repeat.

Greedy Algorithm



- Pick activity you can add with the smallest finish time.
- Repeat.

Greedy Algorithm



- Pick activity you can add with the smallest finish time.
- Repeat.

At least it's fast

- Running time:
 - $O(n)$ if the activities are already sorted by finish time.
 - Otherwise $O(n\log(n))$ if you have to sort them first.

What makes it **greedy**?

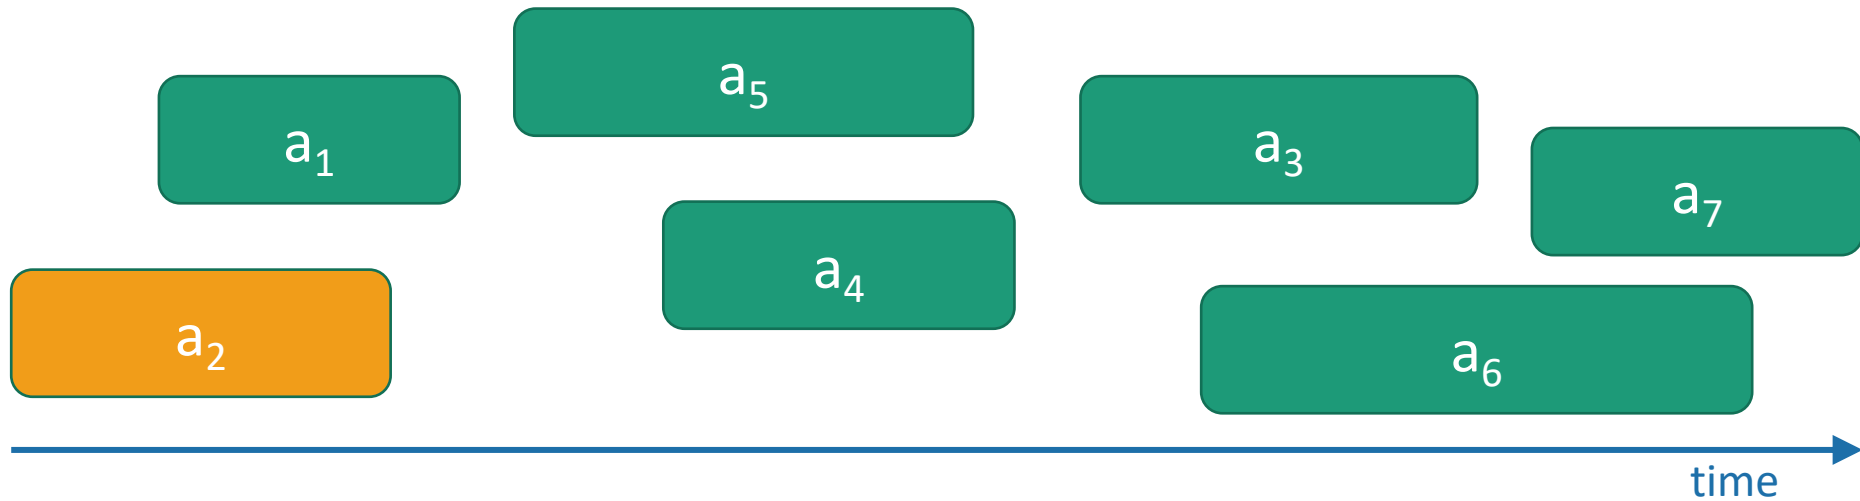
- At each step in the algorithm, make a choice.
 - Hey, I can increase my activity set by one,
 - And leave lots of room for future choices,
 - Let's do that and hope for the best!!!
- **Hope** that at the end of the day, this results in a globally optimal solution.



Three Questions

1. Does this greedy algorithm for activity selection work?
 - Yes. (We will see why in a moment...)
2. In general, when are greedy algorithms a good idea?
 - When the problem exhibits especially nice optimal substructure.
3. The “greedy” approach is often the first you’d think of...
 - Why are we getting to it now, in Week 9?
 - Proving that greedy algorithms work is often not so easy...

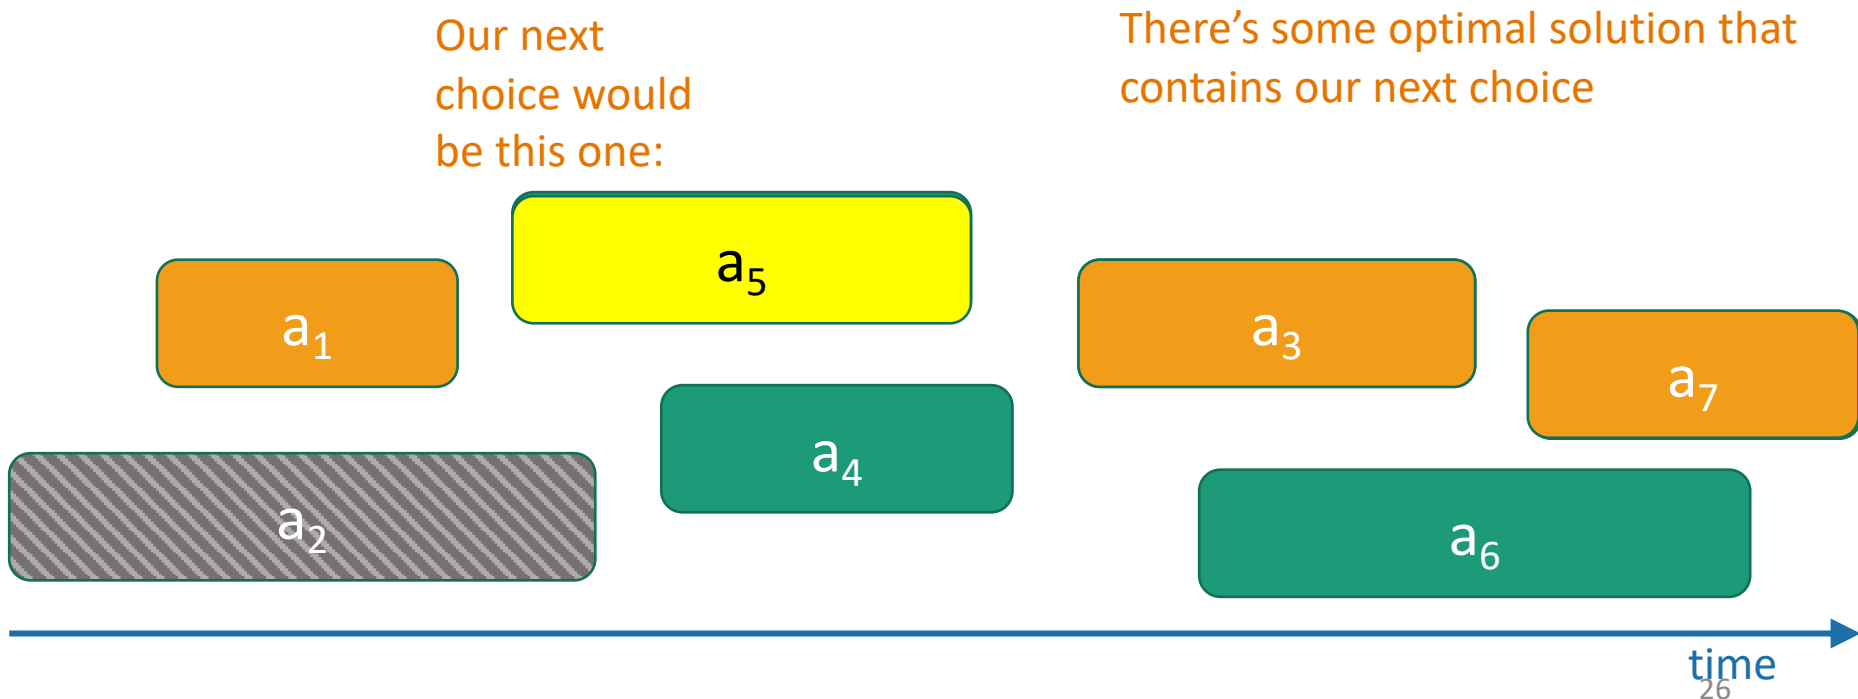
Back to Activity Selection



- Pick activity you can add with the smallest finish time.
- Repeat.

Why does it work?

- Whenever we make a choice, **we don't rule out an optimal solution.**



Assuming we can prove that

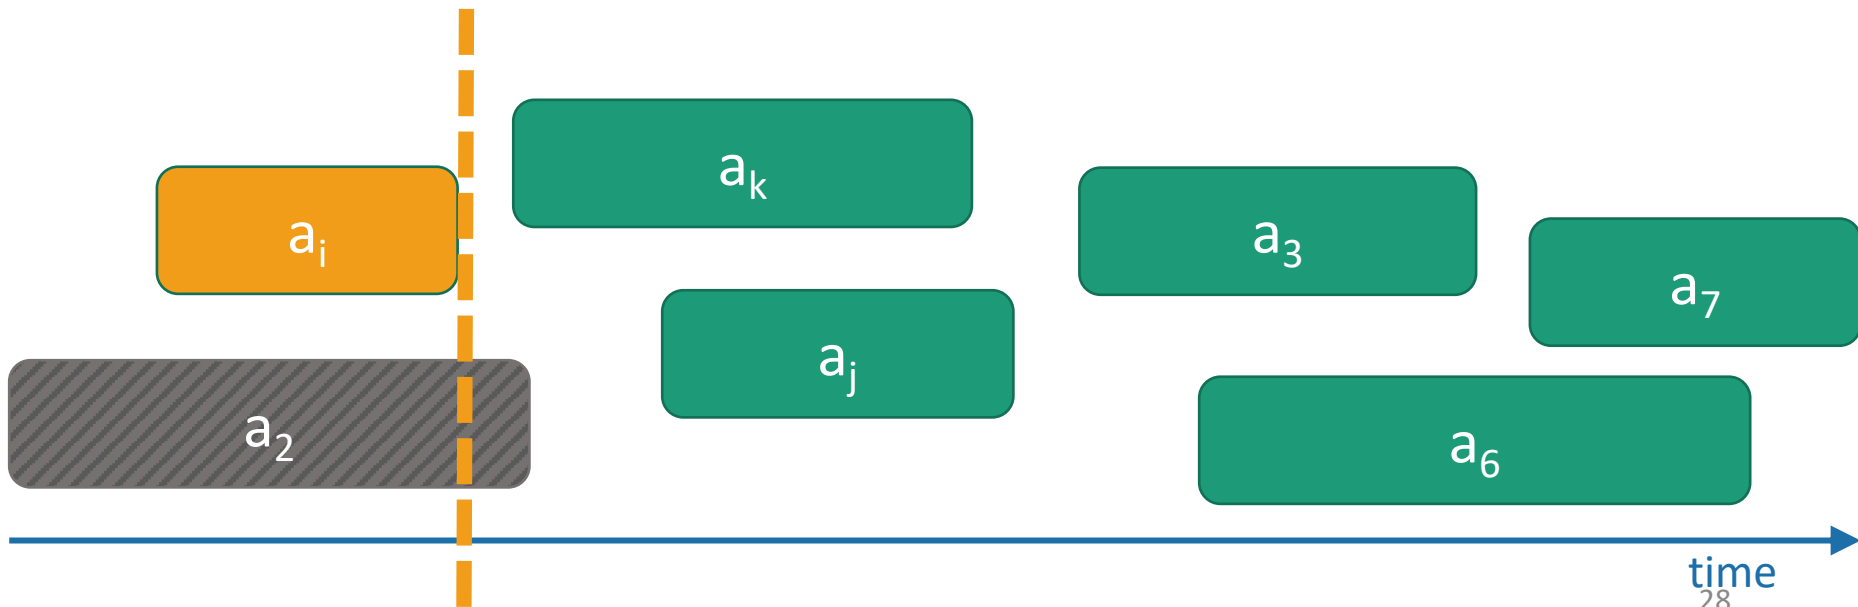
- **We never rule out an optimal solution**
- At the end of the algorithm, we've got some solution.
- So it must be optimal.



Lucky the Lackadaisical Lemur

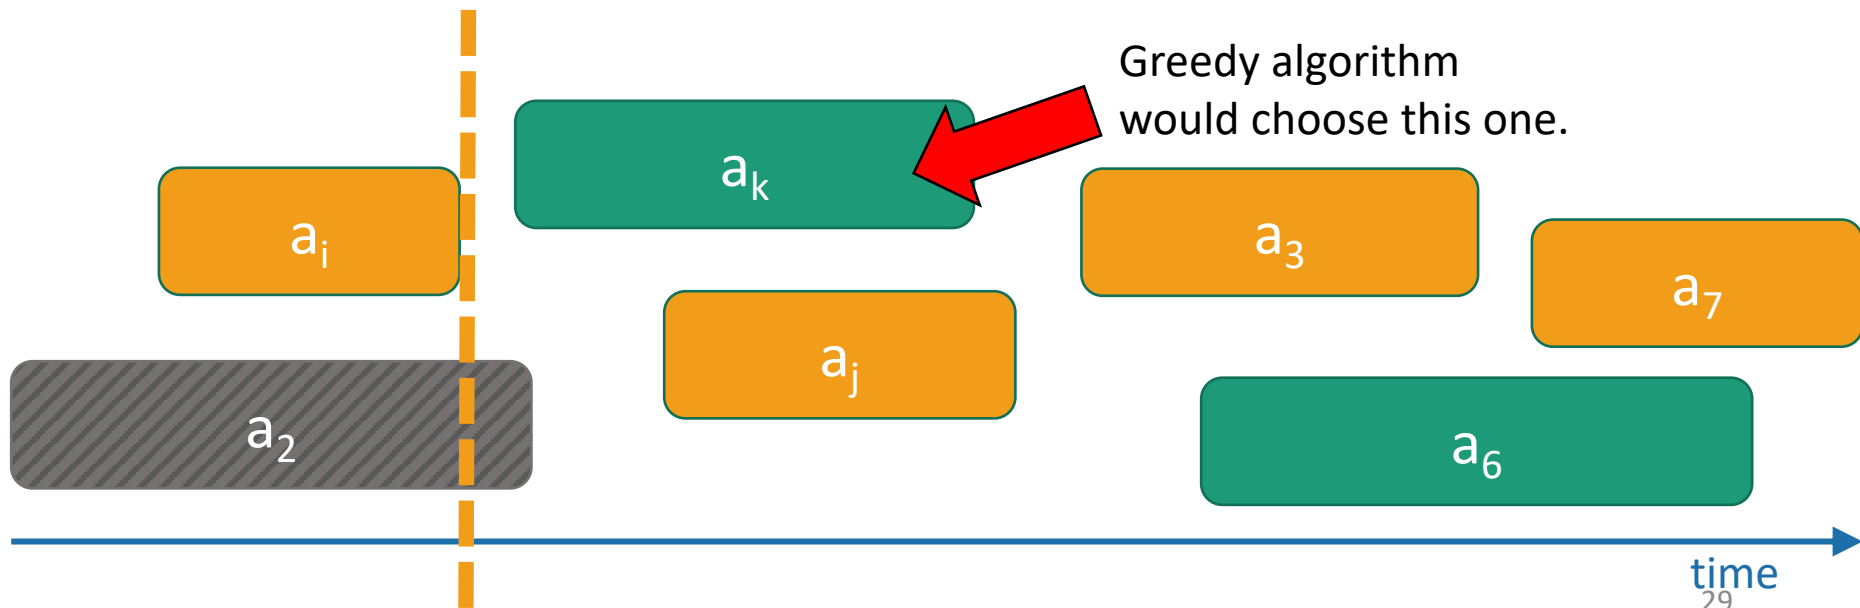
We never rule out an optimal solution

- Suppose we've already chosen a_i , and there is still an optimal solution T^* that extends our choices.



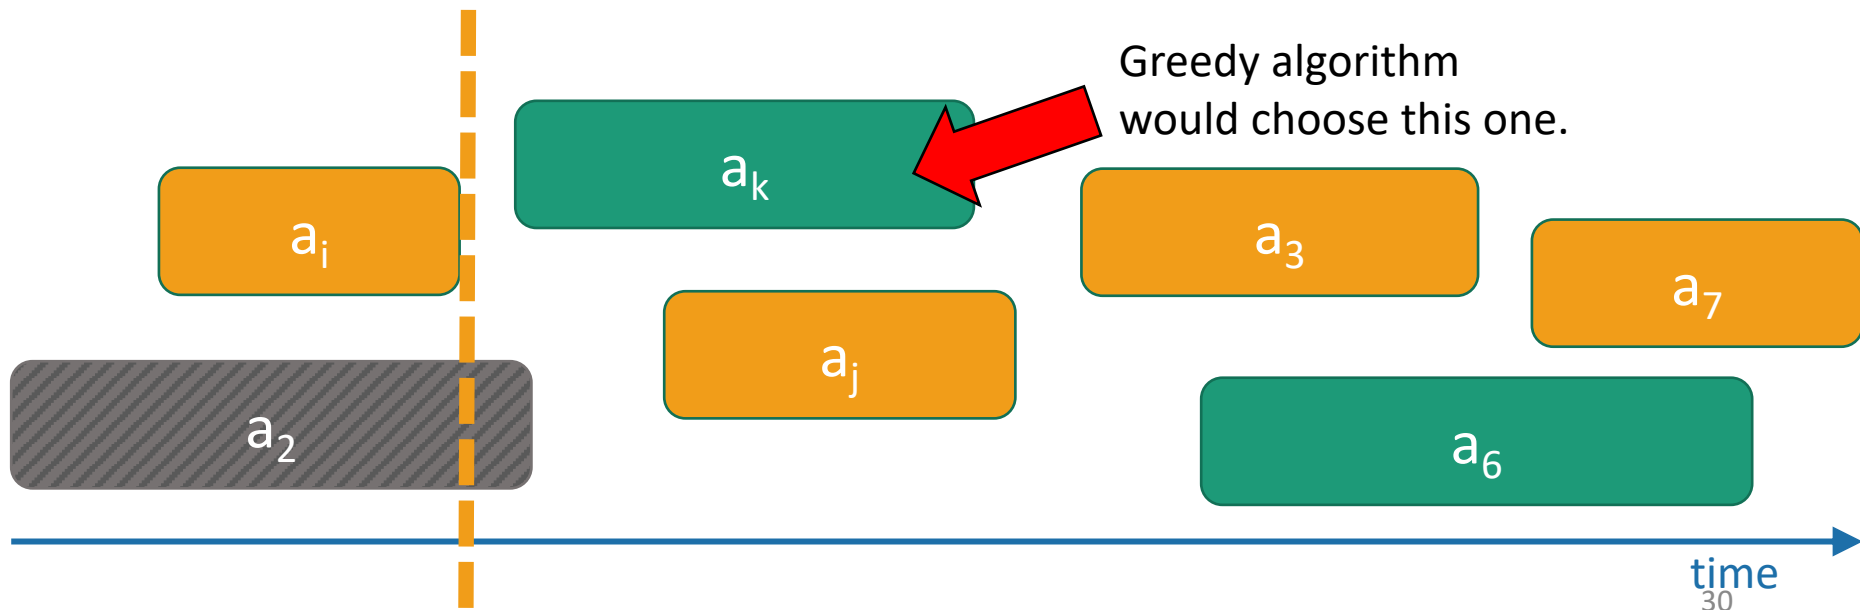
We never rule out an optimal solution

- Suppose we've already chosen a_i , and there is still an optimal solution T^* that extends our choices.
- Now consider the next choice we make, say it's a_k .
- If a_k is in T^* , we're still on track.



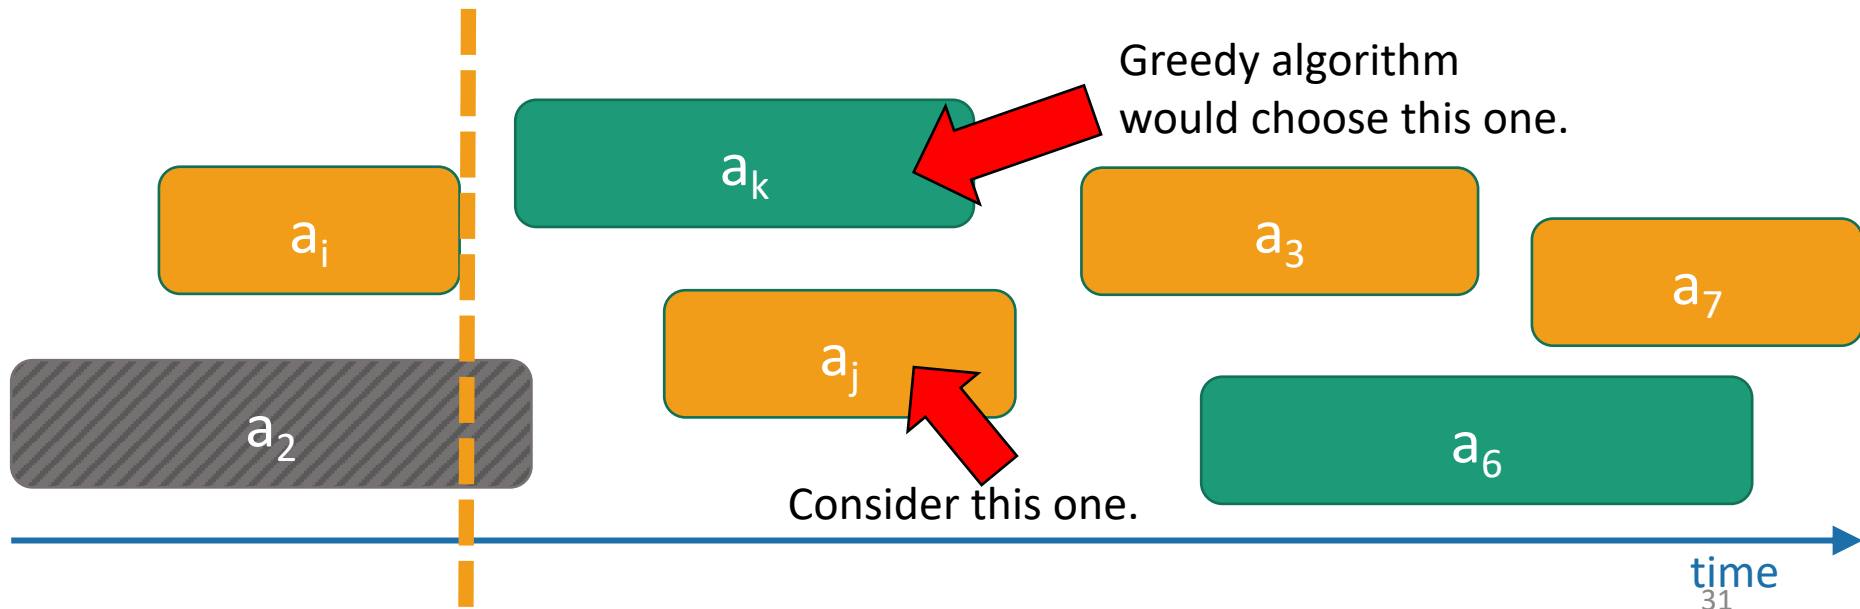
We never rule out an optimal solution

- Suppose we've already chosen a_i , and there is still an optimal solution T^* that extends our choices.
- Now consider the next choice we make, say it's a_k .
- If a_k is **not** in T^* ...



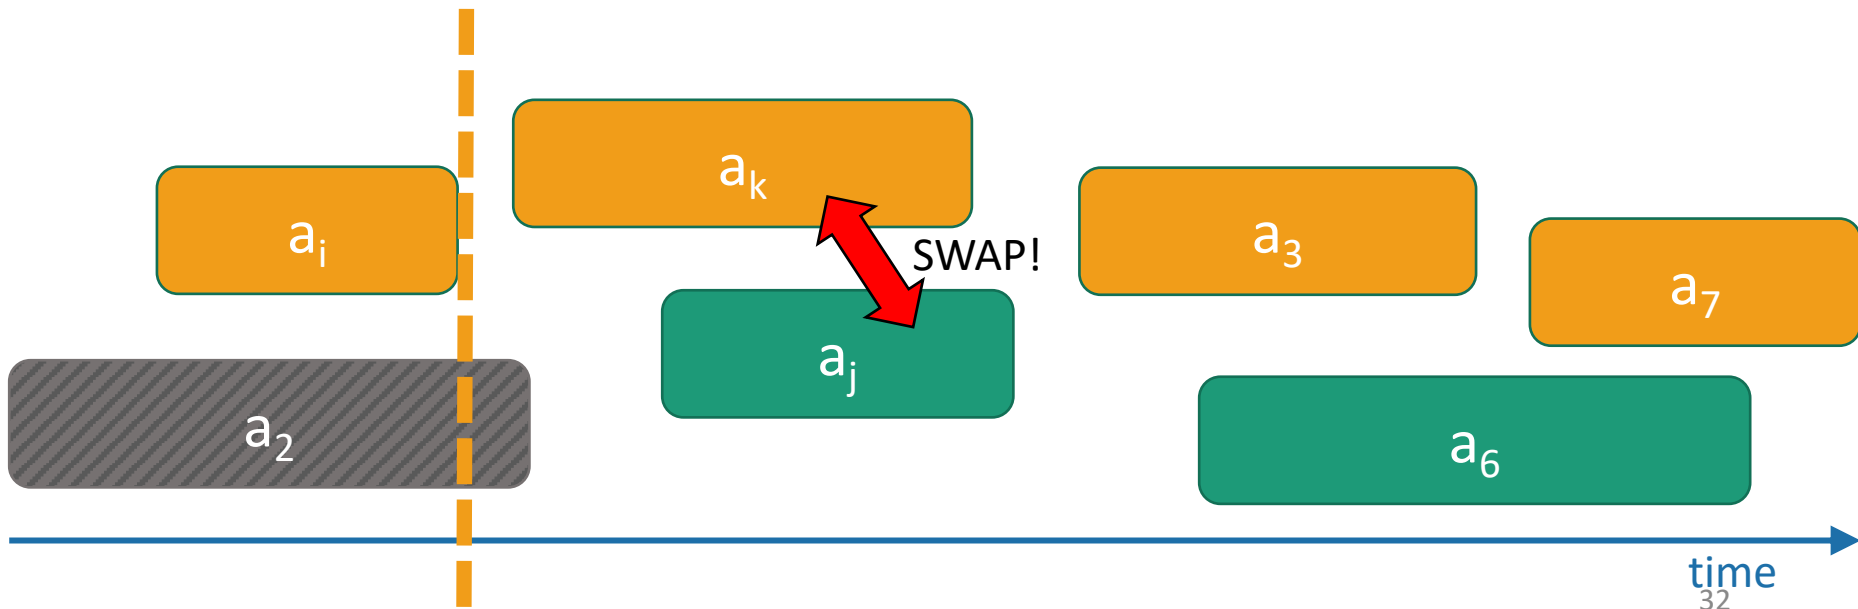
We never rule out an optimal solution ctd.

- If a_k is **not** in T^* ...
- Let a_j be the activity in T^* with the smallest end time.
- Now consider schedule T you get by swapping a_j for a_k



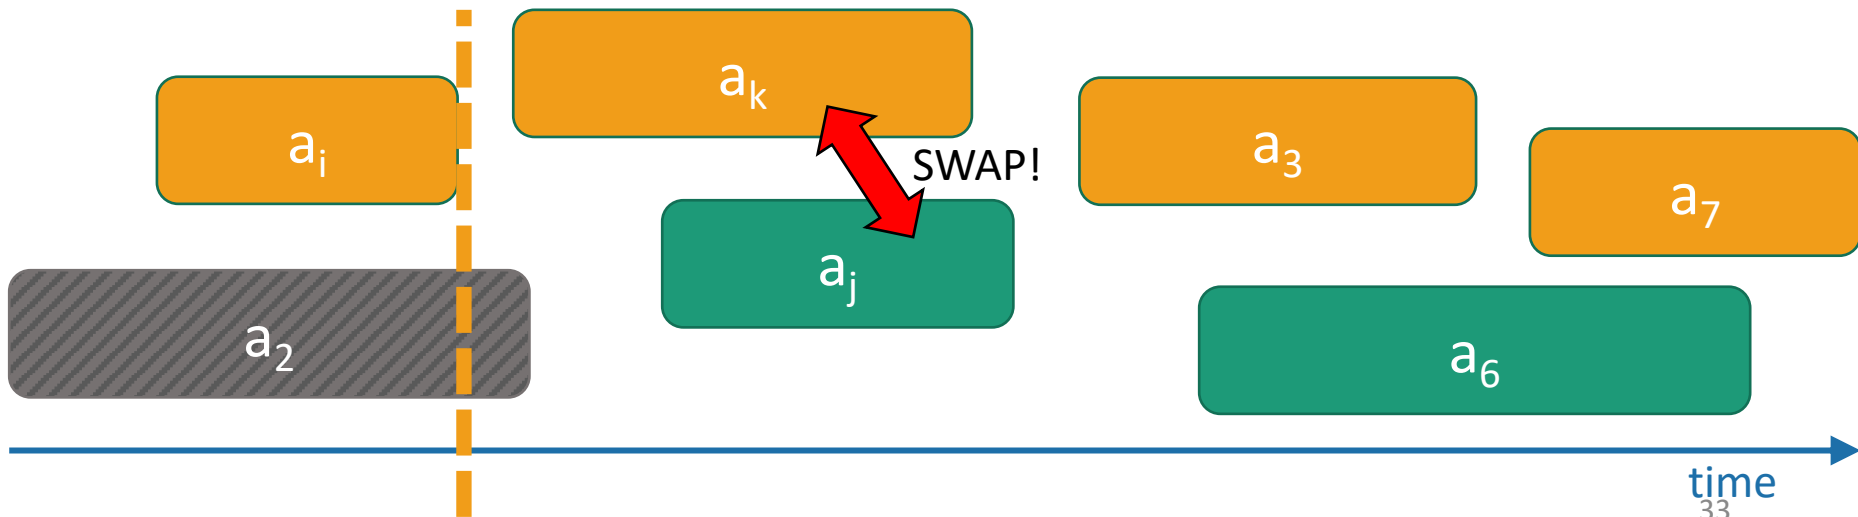
We never rule out an optimal solution cld.

- If a_k is **not** in T^* ...
- Let a_j be the activity in T^* with the smallest end time.
- Now consider schedule T you get by swapping a_j for a_k



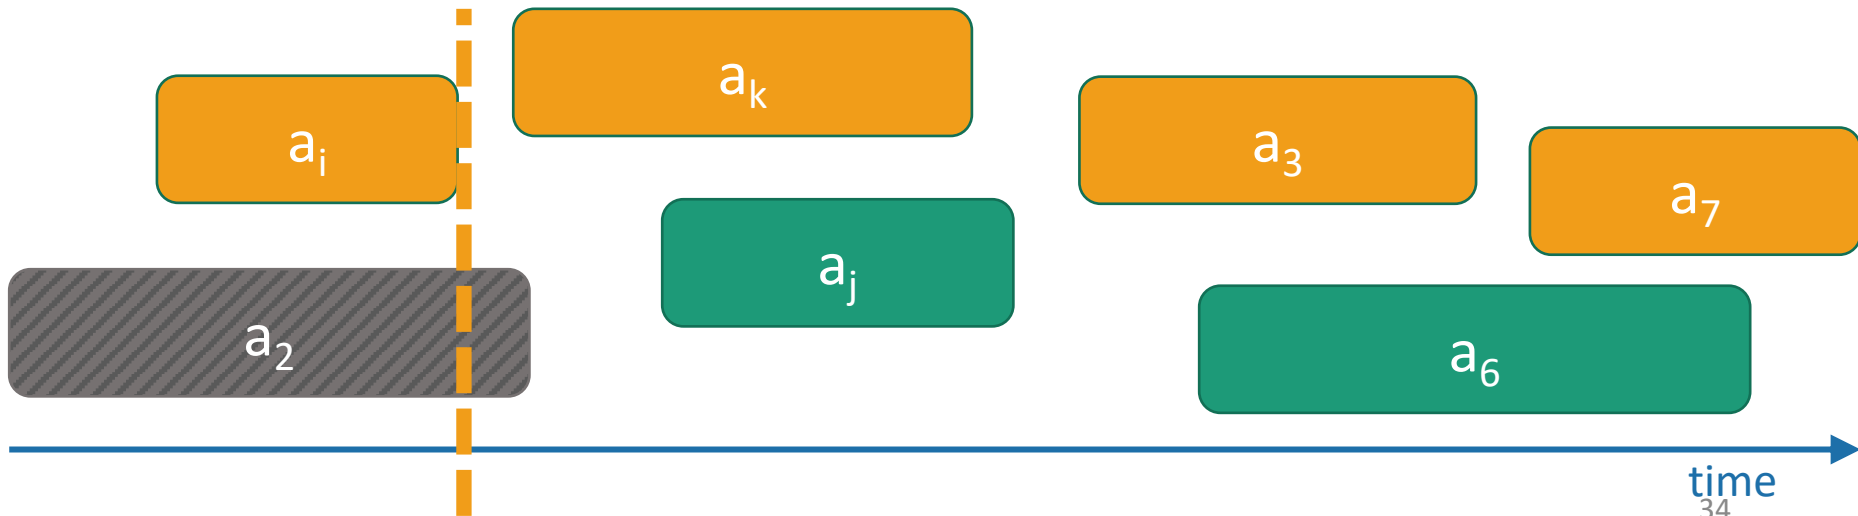
We never rule out an optimal solution cld.

- This schedule T is still allowed.
 - Since a_k has the smallest ending time, it ends before a_j .
 - Thus, a_k doesn't conflict with anything chosen after a_j .
- And, T is still optimal.
 - It has the same number of activities as T^* .



We never rule out an optimal solution_{ctd.}

- We've just shown:
 - If there was an optimal solution that extends the choices we made so far...
 - ...then there is an optimal schedule that also contains our next greedy choice a_k .



So the algorithm is correct

- We never rule out an optimal solution
- At the end of the algorithm, we've got some solution.
- So it must be optimal.



Lucky the Lackadaisical Lemur



So the algorithm is correct



Plucky the Pedantic Penguin

- Inductive Hypothesis:
 - After adding the t 'th thing, there is an optimal solution that extends the current solution.
- Base case:
 - After adding zero activities, there is an optimal solution extending that.
- Inductive step:
 - **We just did that!**
- Conclusion:
 - After adding the last activity, there is an optimal solution that extends the current solution.
 - The current solution is the only solution that extends the current solution.
 - So the current solution is optimal.

Three Questions

1. Does this greedy algorithm for activity selection work?
 - Yes. 
2. In general, when are greedy algorithms a good idea?
 - When the problem exhibits especially nice optimal substructure.
3. The “greedy” approach is often the first you’d think of...
 - Why are we getting to it now, in Week 9?
 - Proving that greedy algorithms work is often not so easy... 

Common strategy for greedy algorithms

- Make a **series of choices**.
- Show that, at each step, our choice **won't rule out an optimal solution** at the end of the day.
- After we've made all our choices, we haven't ruled out an optimal solution, **so we must have found one**.



Common strategy (formally) for greedy algorithms



“Success” here means
“finding an optimal solution.”




- Inductive Hypothesis:
 - After greedy choice t , you haven't ruled out success.
- Base case:
 - Success is possible before you make any choices.
- Inductive step:
 - If you haven't ruled out success after choice t , then you won't rule out success after choice $t+1$.
- Conclusion:
 - If you reach the end of the algorithm and haven't ruled out success then you must have succeeded.

Common strategy

for showing we don't rule out success

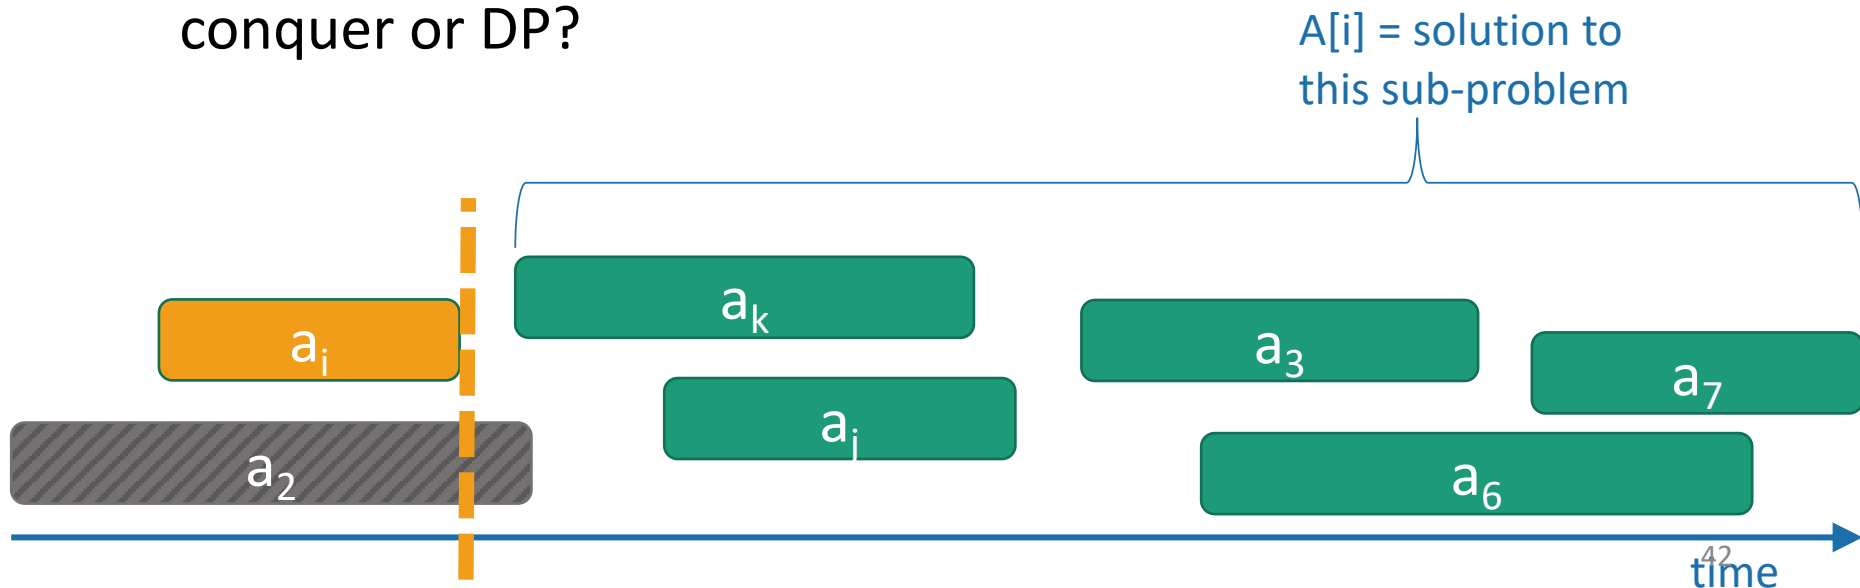
- Suppose that you're on track to make an optimal solution T^* .
 - Eg, after you've picked activity i , you're still on track.
- Suppose that T^* *disagrees* with your next greedy choice.
 - Eg, it *doesn't* involve activity k .
- Manipulate T^* in order to make a solution T that's not worse but that *agrees* with your greedy choice.
 - Eg, swap whatever activity T^* did pick next with activity k .

Three Questions

1. Does this greedy algorithm for activity selection work?
 - Yes. 
2. In general, when are greedy algorithms a good idea?
 - When the problem exhibits especially nice optimal substructure. 
3. The “greedy” approach is often the first you’d think of...
 - Why are we getting to it now, in Week 9?
 - Proving that greedy algorithms work is often not so easy... 

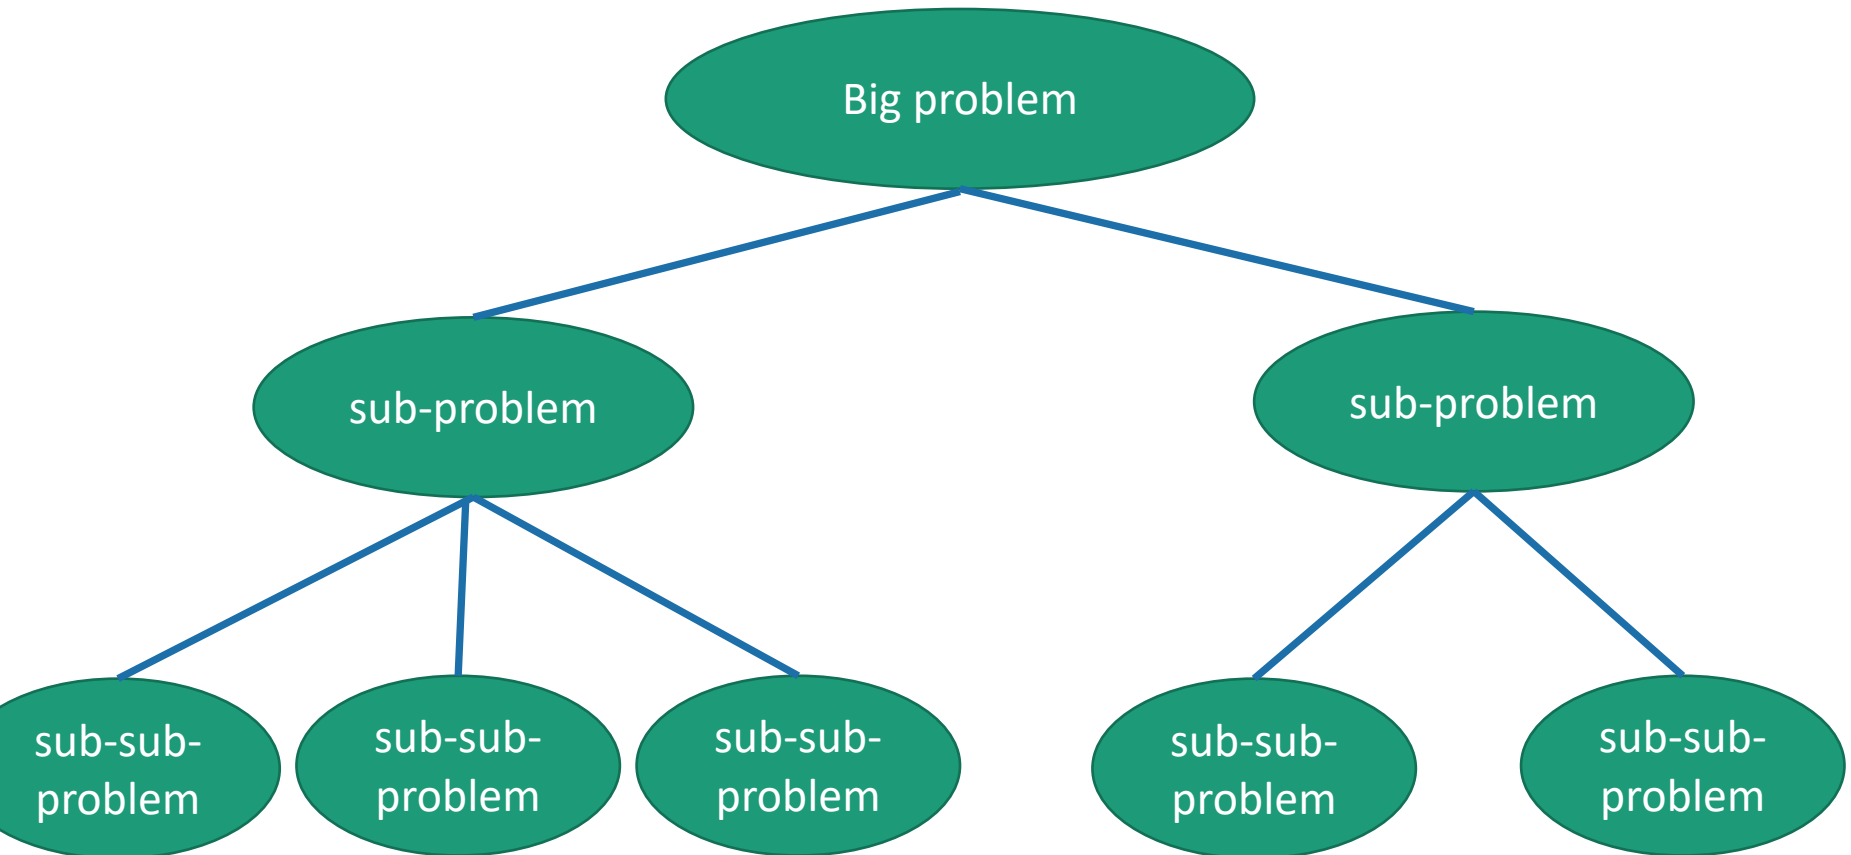
Optimal sub-structure in greedy algorithms

- Our greedy activity selection algorithm exploited a natural sub-problem structure:
 $A[i]$ = number of activities you can do after the end of activity i
- How does this substructure relate to that of divide-and-conquer or DP?




Sub-problem graph view

- Divide-and-conquer:




Sub-problem graph view

- Dynamic Programming:




Sub-problem graph view

- Greedy algorithms:



Sub-problem graph view

- Greedy algorithms:






- Not only is there **optimal sub-structure**:
 - optimal solutions to a problem are made up from optimal solutions of sub-problems
- but each problem **depends on only one sub-problem**.

Write a DP version of activity selection (where you fill in a table)! [See hidden slides in the .pptx file for one way]



Three Questions

1. Does this greedy algorithm for activity selection work?
 - Yes. 
2. In general, when are greedy algorithms a good idea?
 - When they exhibit especially nice optimal substructure. 
3. The “greedy” approach is often the first you’d think of...
 - Why are we getting to it now, in Week 9?
 - Proving that greedy algorithms work is often not so easy. 

Let's see a few more examples

Another example: Scheduling

CS161 HW

Personal Hygiene

Math HW

Administrative stuff for student club

Econ HW

Do laundry

Meditate

Practice musical instrument

Read CLRS


Have a social life

Sleep



Scheduling

- n tasks
- Task i takes t_i hours
- For every hour that passes until task i is done, pay c_i



- CS161 HW, then Sleep: costs $10 \cdot 2 + (10 + 8) \cdot 3 = 74$ units
- Sleep, then CS161 HW: costs $8 \cdot 3 + (10 + 8) \cdot 2 = 60$ units

Optimal substructure


- This problem breaks up nicely into sub-problems:

Suppose this is the optimal schedule:



Then this must be the optimal schedule on just jobs B,C,D.

Why?



Think-pair-share

Optimal substructure

- Seems amenable to a greedy algorithm:

Take the best job first

Then solve this problem



Take the best job first

Then solve this problem



Take the best job first

Then solve this problem



(That one's easy 😊) 66

What does “best” mean?

AB is better than **BA** when:


$$xz + (x + y)w \leq yw + (x + y)z$$

$$xz + xw + yw \leq yw + xz + yz$$

$$wx \leq yz$$

$$\frac{w}{y} \leq \frac{z}{x}$$

- Of these two jobs, which should we do first?



Cost: z units per hour until it's done.

Cost: w units per hour until it's done.

- Cost(**A then B**) = $x \cdot z + (x + y) \cdot w$
- Cost(**B then A**) = $y \cdot w + (x + y) \cdot z$

What matters is the ratio:

cost of delay
time it takes

“Best” means
biggest ratio.⁶⁷

Idea for greedy algorithm

- Choose the job with the biggest $\frac{\text{cost of delay}}{\text{time it takes}}$ ratio.

Lemma

This greedy choice doesn't rule out success

- Suppose you have already chosen some jobs, and haven't yet ruled out success:


Already
chosen E

There's some way to order
A, B, C, D that's optimal...



- Then if you choose the next job to be the one left that maximizes the ratio **cost/time**, you still won't rule out success.
- Proof sketch:**
 - Say Job B maximizes this ratio, but it's not the next job in the opt. soln.

How can we manipulate the optimal solution above to make an optimal solution where B is the next job we choose?



Lemma

This greedy choice doesn't rule out success

- Suppose you have already chosen some jobs, and haven't yet ruled out success:

Already
chosen E



There's some way to order
A, B, C, D that's optimal...

- Then if you choose the next job to be the one left that maximizes the ratio **cost/time**, you still won't rule out success.

- Proof sketch:**

- Say Job B maximizes this ratio, but it's not the next job in the opt. soln.
- Switch A and B! Nothing else will change, and we just showed that the cost of the solution won't increase.



- Repeat until B is first.




- Now this is an optimal schedule where B is first.

Back to our framework for proving correctness of greedy algorithms

- Inductive Hypothesis:
 - After greedy choice t , you haven't ruled out success.
- Base case:
 - Success is possible before you make any choices.
- Inductive step:
 - If you haven't ruled out success after choice t , then you won't rule out success after choice $t+1$.
- Conclusion:
 - If you reach the end of the algorithm and haven't ruled out success then you must have succeeded.

Just did the inductive step!



Fill in the details!



Greedy Scheduling Solution

- **scheduleJobs(JOBS):**
 - Sort JOBS in decreasing order by the ratio:
 - $r_i = \frac{c_i}{t_i} = \frac{\text{cost of delaying job } i}{\text{time job } i \text{ takes to complete}}$
 - **Return JOBS**

Running time: $O(n \log(n))$



Now you can go about your schedule peacefully, in the optimal way.

What have we learned?

- A **greedy algorithm** works for scheduling
- This followed the same outline as the previous example:
 - Identify **optimal substructure**:



- Find a way to make choices that **won't rule out an optimal solution**.
 - largest cost/time ratios first.

One more example

Huffman coding

- everyday english sentence

- 01100101 01110110 01100101 01110010 01111001 01100100 01100001
01111001 00100000 01100101 01101110 01100111 01101100 01101001
01110011 01101000 00100000 01110011 01100101 01101110 01110100
01100101 01101110 01100011 01100101

- qwertyui_opasdfg+hjklzxcv

- 01110001 01110111 01100101 01110010 01110100 01111001 01110101
01101001 01011111 01101111 01110000 01100001 01110011 01100100
01100110 01100111 00101011 01101000 01101010 01101011 01101100
01111010 01111000 01100011 01110110

One more example

Huffman coding

ASCII is pretty wasteful for English sentences. If **e** shows up so often, we should have a more parsimonious way of representing it!


- **e**veryday **e**nglish **s**entence

- **01100101** 01110110 **01100101** 01110010 01111001 01100100 01100001
01111001 00100000 **01100101** 01101110 01100111 01101100 01101001
01110011 01101000 00100000 01110011 **01100101** 01101110 01110100
01100101 01101110 01100011 **01100101**

- qwertyui_opasdfg+hjklzxcv


- 01110001 01110111 01100101 01110010 01110100 01111001 01110101
01101001 01011111 01101111 01110000 01100001 01110011 01100100
01100110 01100111 00101011 01101000 01101010 01101011 01101100
01111010 01111000 01100011 01110110

Suppose we have some distribution on characters



Suppose we have some distribution on characters

For simplicity,
let's go with this
made-up example




Try 0

(like ASCII)

- Every letter is assigned a **binary string** of three bits.


Wasteful!

- 110 and 111 are never used.
- We should have a shorter way of representing A.



Try 1


- Every letter is assigned a **binary string** of one or two bits.
- The more frequent letters get the shorter strings.
- **Problem:**
 - Does 000 mean AAA or BA or AB?



Confusingly, “prefix-free codes” are also sometimes called “prefix codes” (including in CLRS).

Try 2: prefix-free coding


- Every letter is assigned a **binary string**.
- More frequent letters get shorter strings.
- No encoded string is a **prefix** of any other.



Confusingly, “prefix-free codes” are also sometimes called “prefix codes” (including in CLRS).

Try 2: prefix-free coding


- Every letter is assigned a **binary string**.
- More frequent letters get shorter strings.
- No encoded string is a **prefix** of any other.



Confusingly, “prefix-free codes” are also sometimes called “prefix codes” (including in CLRS).

Try 2: prefix-free coding


- Every letter is assigned a **binary string**.
- More frequent letters get shorter strings.
- No encoded string is a **prefix** of any other.



Confusingly, “prefix-free codes” are also sometimes called “prefix codes” (including in CLRS).

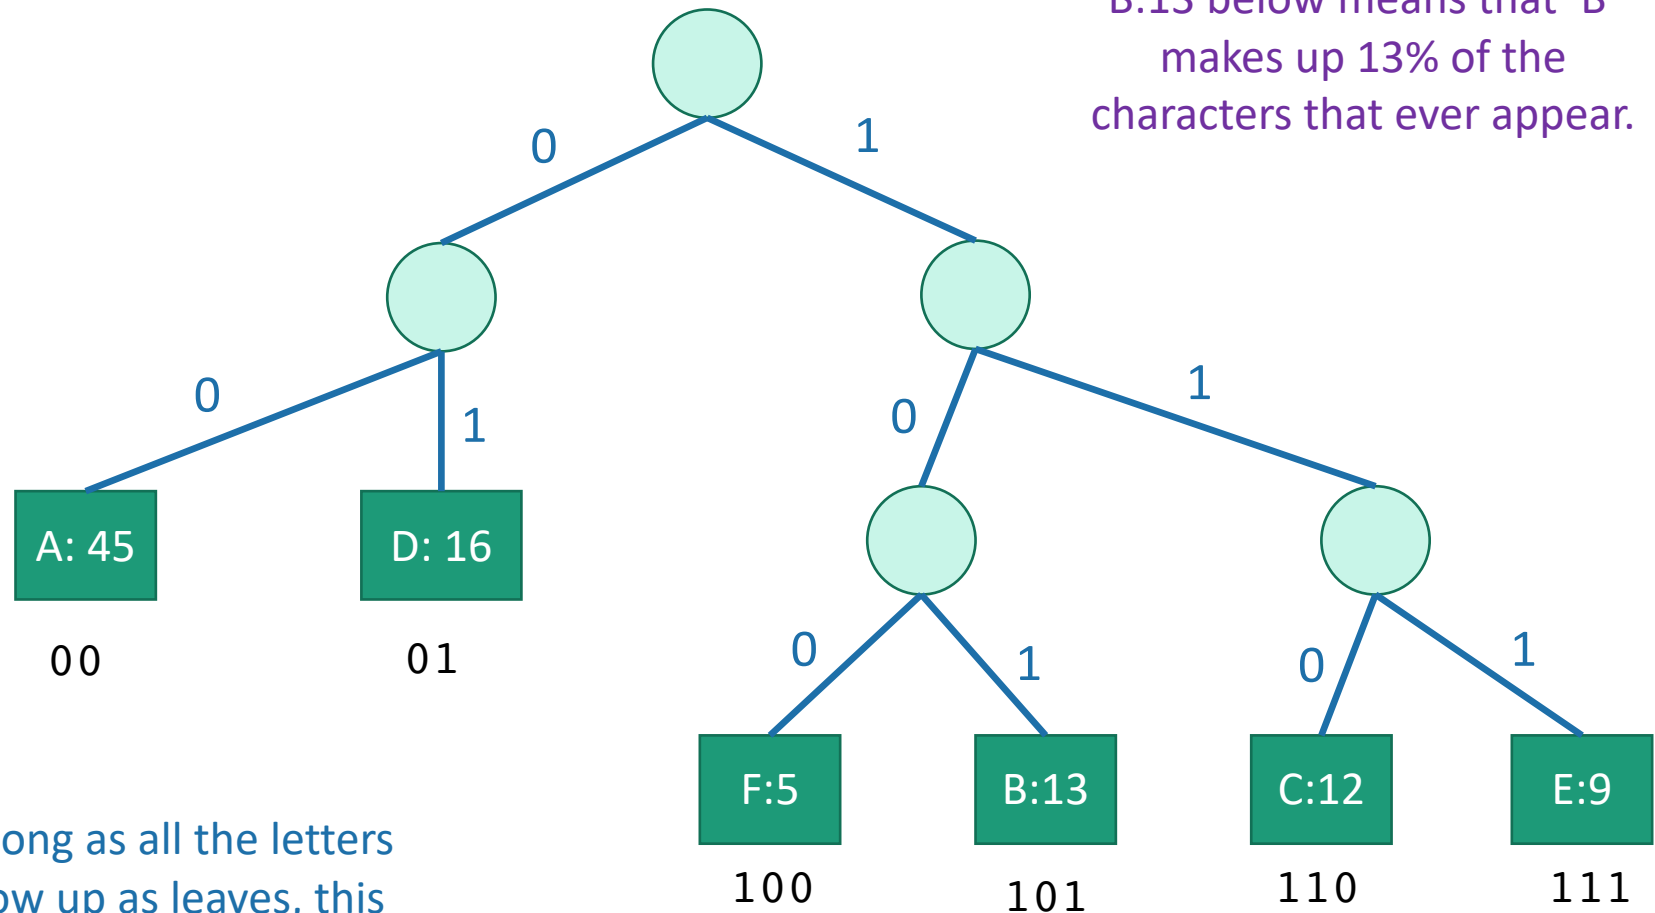
Try 2: prefix-free coding

- Every letter is assigned a **binary string**.
- More frequent letters get shorter strings.
- No encoded string is a **prefix** of any other.



A prefix-free code is a tree

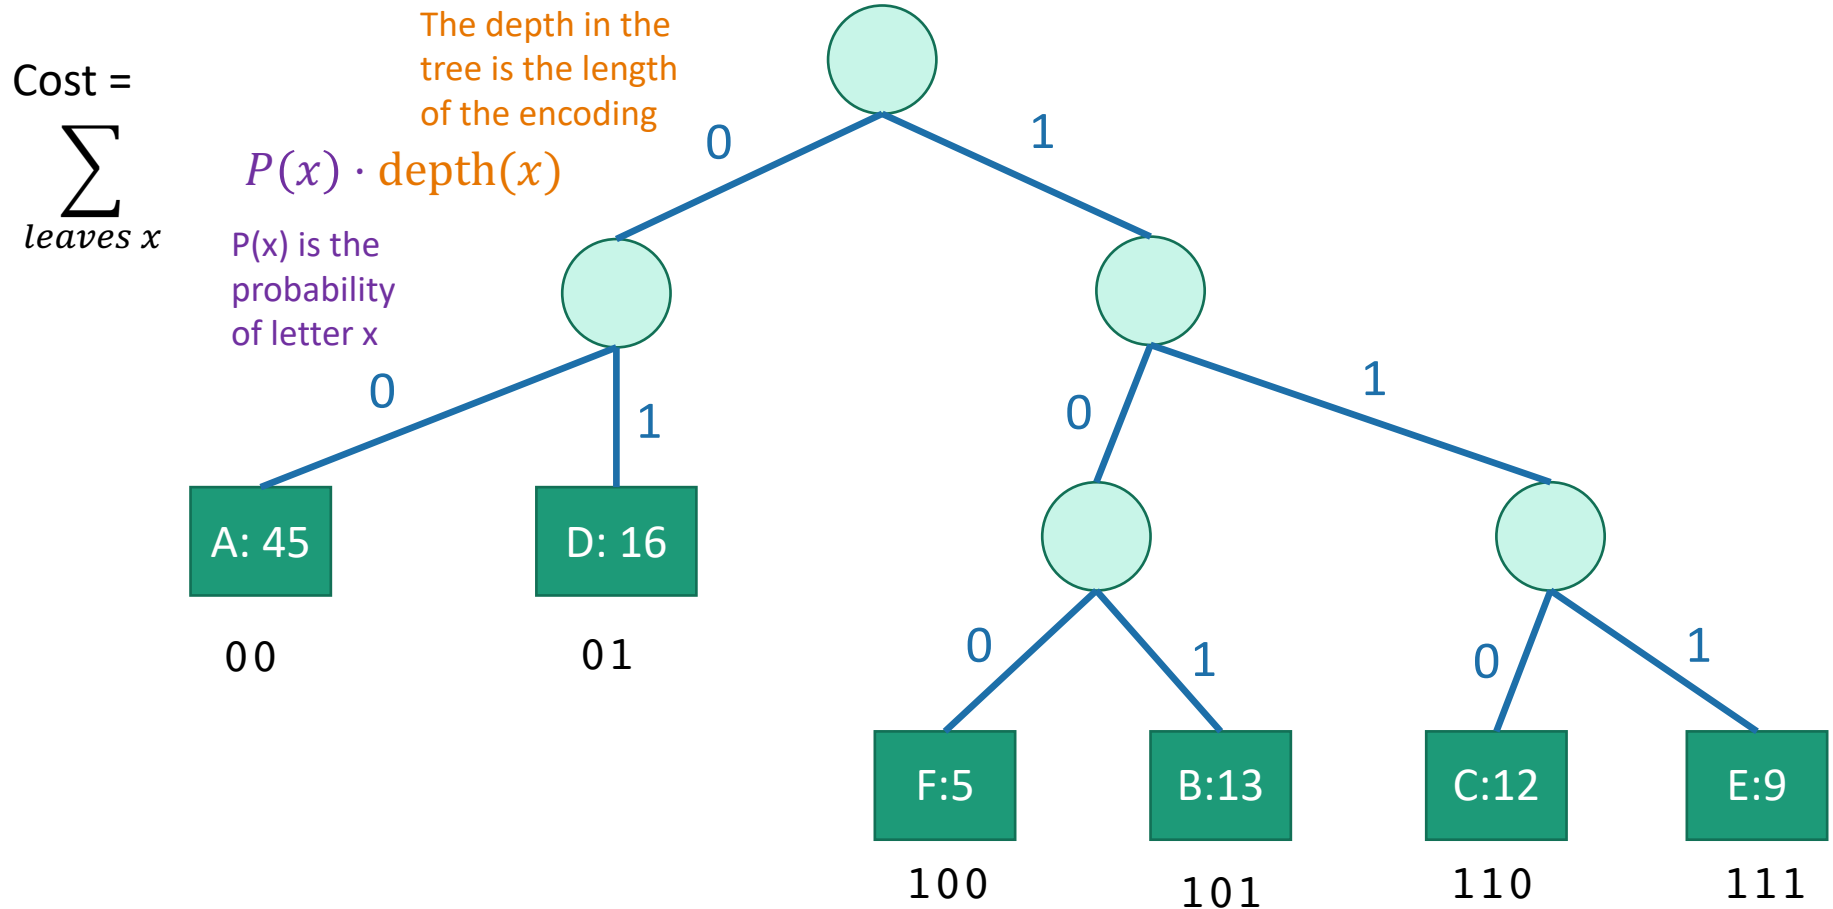
B:13 below means that 'B'
makes up 13% of the
characters that ever appear.



As long as all the letters
show up as leaves, this
code is **prefix-free**.

How good is a tree?

- Imagine choosing a letter at random from the language.
 - Not uniform, but according to our histogram!
- The **cost of a tree** is the expected length of the encoding of that letter.



Expected cost of encoding a letter with this tree:

$$2(0.45 + 0.16) + 3(0.05 + 0.13 + 0.12 + 0.09) = 2.39$$

Question

- Given a distribution P on letters, find the lowest-cost tree, where

$$\text{cost}(\text{tree}) = \sum_{\text{leaves } x} P(x) \cdot \text{depth}(x)$$

$P(x)$ is the probability of letter x


The depth in the tree is the length of the encoding

Greedy algorithm

- Greedily build sub-trees from the bottom up.
- Greedy goal: less frequent letters should be further down the tree.


Solution

greedily build subtrees, starting with the infrequent letters




Solution

greedily build subtrees, starting with the infrequent letters




Solution

greedily build subtrees, starting with the infrequent letters



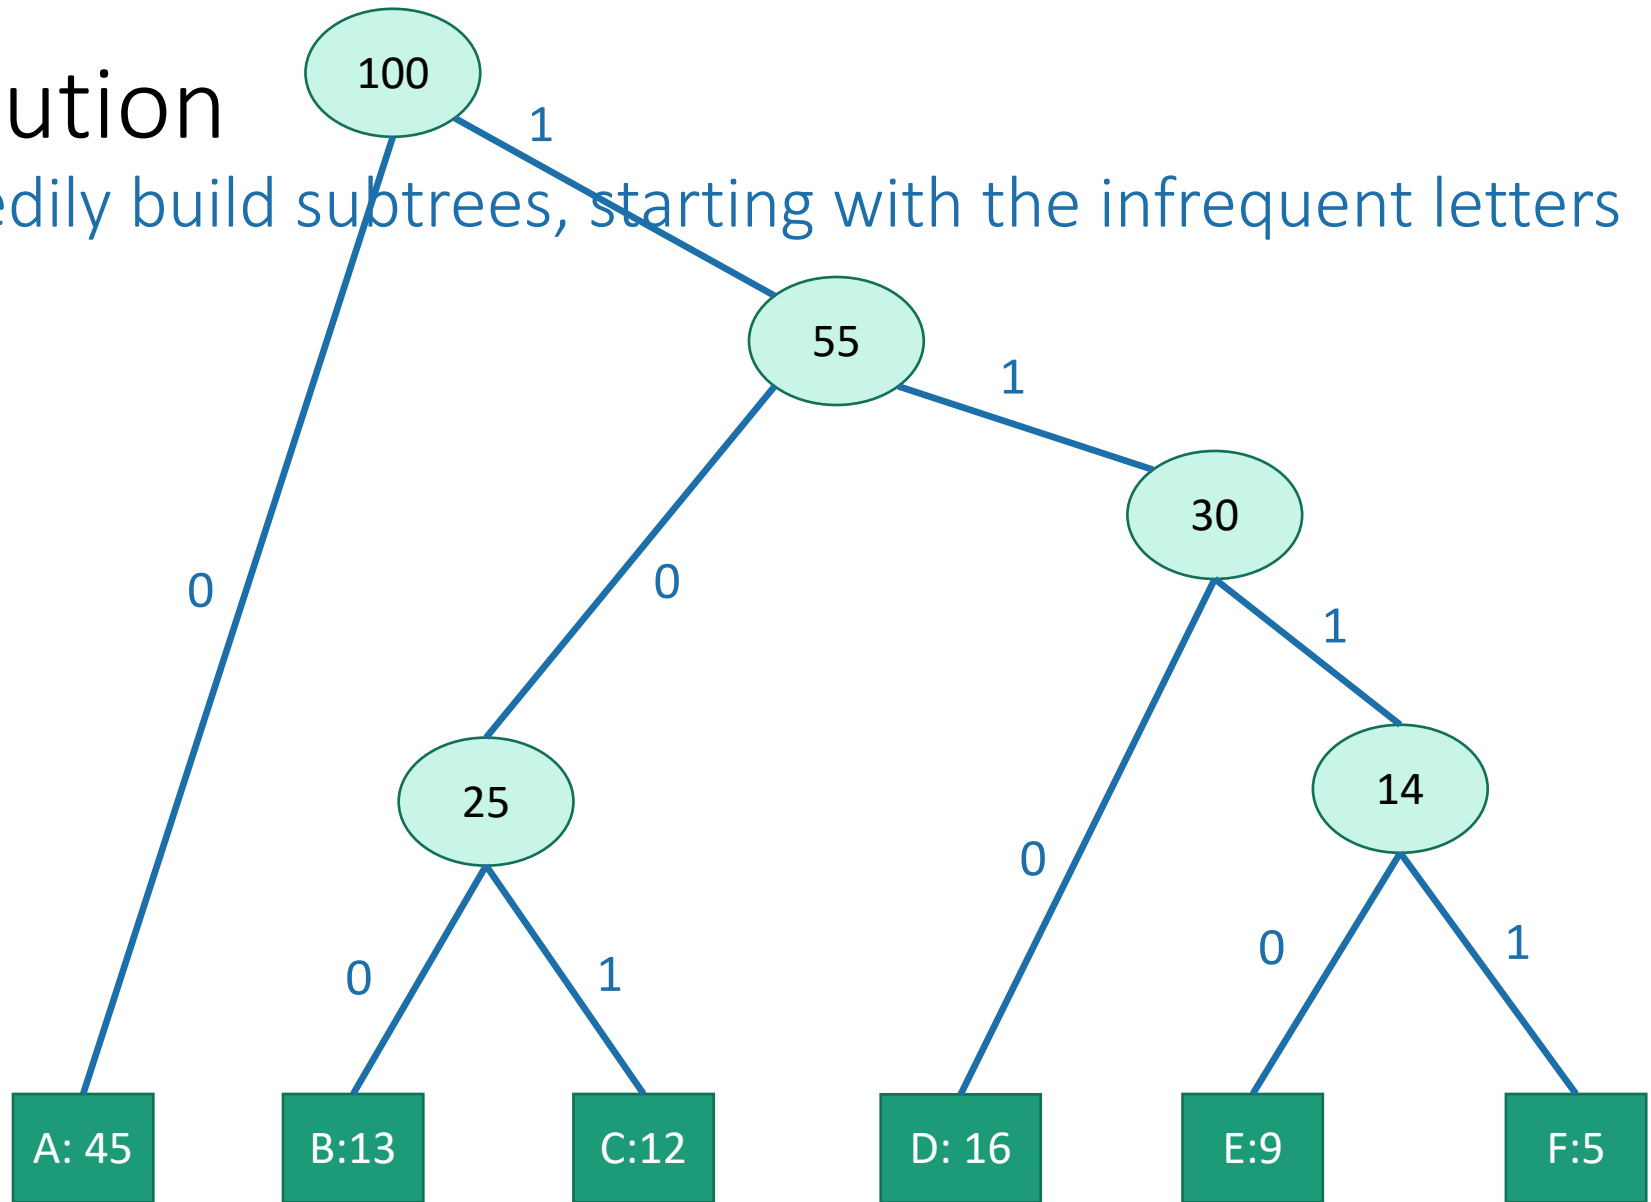
Solution

greedily build subtrees, starting with the infrequent letters



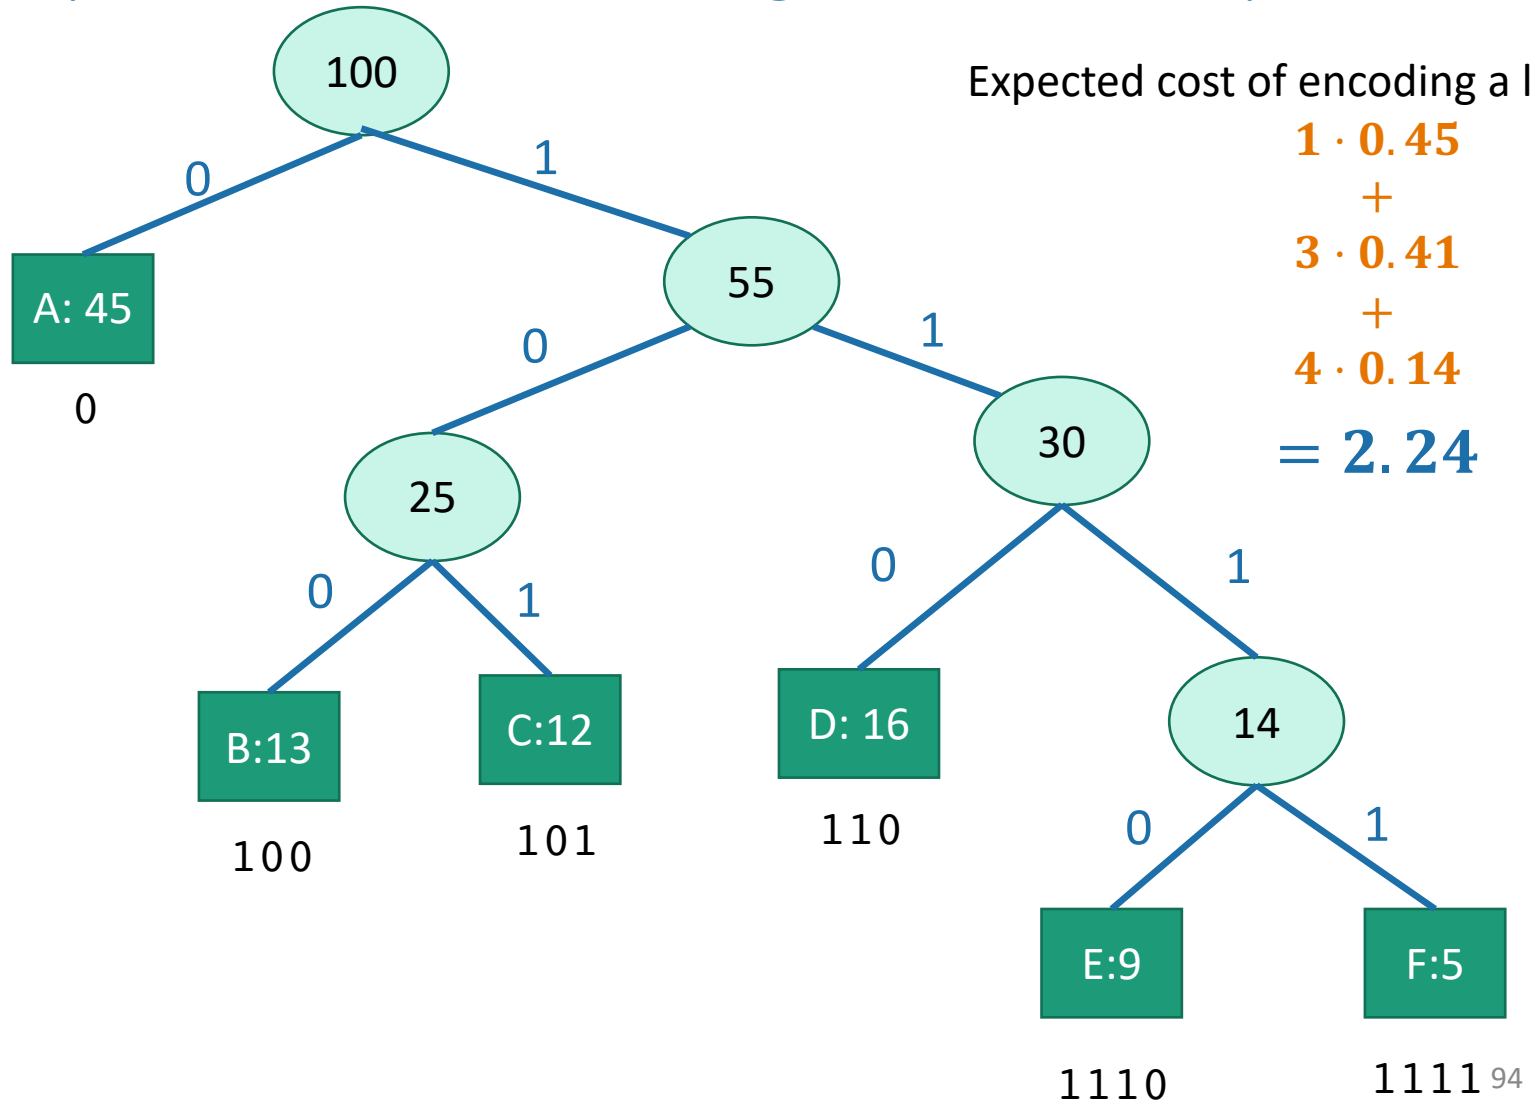
Solution

greedily build subtrees, starting with the infrequent letters



Solution

greedily build subtrees, starting with the infrequent letters



What exactly was the algorithm?


- Create a node like **D: 16** for each letter/frequency
 - The key is the frequency (16 in this case)
- Let **CURRENT** be the list of all these nodes.
- **while** len(**CURRENT**) > 1:
 - **X** and **Y** ← the nodes in **CURRENT** with the smallest keys.
 - Create a new node **Z** with **Z.key = X.key + Y.key**
 - Set **Z.left = X, Z.right = Y**
 - Add **Z** to **CURRENT** and remove **X** and **Y**
- return **CURRENT**[0]

A: 45

B: 13

C: 12

D: 16



This is called Huffman Coding:

- Create a node like D: 16 for each letter/frequency
 - The key is the frequency (16 in this case)
- Let **CURRENT** be the list of all these nodes.
- **while** len(**CURRENT**) > 1:
 - **X** and **Y** ← the nodes in **CURRENT** with the smallest keys.
 - Create a new node **Z** with **Z.key = X.key + Y.key**
 - Set **Z.left = X**, **Z.right = Y**
 - Add **Z** to **CURRENT** and remove **X** and **Y**
- return **CURRENT**[0]



Does it work?


- Yes.
- We will ***sketch*** a proof here.
- Same strategy:
 - Show that at each step, the choices we are making **won't rule out** an optimal solution.
 - Lemma:
 - Suppose that x and y are the two least-frequent letters. Then there is an optimal tree where x and y are siblings.



Lemma proof idea

If x and y are the two least-frequent letters, there is an optimal tree where x and y are siblings.

- Say that an optimal tree looks like this:




Lowest-level sibling nodes: at least one of them is neither x nor y

- What happens to the cost if we swap x for a ?
 - the cost can't increase; a was more frequent than x , and we just made a 's encoding shorter and x 's longer.
- Repeat this logic until we get an optimal tree with x and y as siblings.
 - The cost never increased so this tree is still optimal.

Lemma proof idea

If x and y are the two least-frequent letters, there is an optimal tree where x and y are siblings.

- Say that an optimal tree looks like this:



Lowest-level sibling nodes: at least one of them is neither x nor y


- What happens to the cost if we swap x for a ?
 - the cost can't increase; a was more frequent than x , and we just made a 's encoding shorter and x 's longer.
- Repeat this logic until we get an optimal tree with x and y as siblings.
 - The cost never increased so this tree is still optimal.

Proof strategy

just like before

- Show that at each step, the choices we are making **won't rule out** an optimal solution.
- Lemma:
 - Suppose that x and y are the two least-frequent letters. Then there is an optimal tree where x and y are siblings.

That's enough to show that we don't rule out optimality after the first step.





Proof strategy

just like before

- Show that at each step, the choices we are making **won't rule out** an optimal solution.
- Lemma:
 - Suppose that x and y are the two least-frequent letters. Then there is an optimal tree where x and y are siblings.



That's enough to show that we don't rule out optimality after the first step.

What about once we start grouping stuff?



Lemma 2

this distinction doesn't really matter



The first thing is an optimal tree on $\{A, B, C, D, E, F\}$

if and only if

the second thing is an optimal tree on $\{A, G, H\}$

Lemma 2

this distinction doesn't really matter

- For a proof:
 - See CLRS, Lemma 16.3
 - Rigorous although presented in a slightly different way
 - See the (optional) Lecture Notes
 - A bit sketchier, but presented in the same way as here
 - Prove it yourself!
 - This is the best!

Getting all the details
isn't that important, but
you should convince
yourself that this is true.



Together

- Lemma 1:
 - Suppose that x and y are the two least-frequent letters.
Then there is an optimal tree where x and y are siblings.
- Lemma 2:
 - We may as well imagine that **CURRENT** contains only leaves.
- These imply:
 - At each step, our choice doesn't rule out an optimal tree.

Write this out formally as a proof by induction! (See skipped slides for a starting point).






The whole argument

After the t' th step, we've got a bunch of current sub-trees:

- Inductive hypothesis:

- after the t' th step,
 - there is an optimal tree containing the current subtrees as “leaves”



- Base case:

- after the 0'th step,
 - there is an optimal tree containing all the characters.


Inductive hyp. asserts that our subtrees can be assembled into an optimal tree:

- Inductive step:


- **TO DO**

- Conclusion:

- after the last step,
 - there is an optimal tree containing this whole tree as a subtree.
- aka,
 - after the last step the tree we've constructed is optimal.



We've got a bunch of current sub-trees:




say that x and y are the two smallest.

Inductive step

- Suppose that the inductive hypothesis holds for $t-1$
 - After $t-1$ steps, there is an optimal tree containing all the current sub-trees as “leaves.”
- Want to show:
 - After t steps, there is an optimal tree containing all the current sub-trees as leaves.


We've got a bunch of current sub-trees:



say that x and y are the two smallest.

Inductive step


- Suppose that the inductive hypothesis holds for $t-1$
 - After $t-1$ steps, there is an optimal tree containing all the current sub-trees as “leaves.”



- By Lemma 2, may as well treat




We've got a bunch of current sub-trees:



say that x and y are the two smallest.

Inductive step

- Suppose that the inductive hypothesis holds for $t-1$
 - After $t-1$ steps, there is an optimal tree containing all the current sub-trees as “leaves.”




- By Lemma 2, may as well treat



- In particular, optimal trees on this new alphabet correspond to optimal trees on the original alphabet.


We've got a bunch of current sub-trees:




say that x and y are the two smallest.

Inductive step


- Suppose that the inductive hypothesis holds for $t-1$
 - After $t-1$ steps, there is an optimal tree containing all the current sub-trees as “leaves.”



- Our algorithm would do this at level t :




We've got a bunch of current sub-trees:



say that x and y are the two smallest.


Inductive step

- Suppose that the inductive hypothesis holds for $t-1$
 - After $t-1$ steps, there is an optimal tree containing all the current sub-trees as “leaves.”




Lemma 1 implies that there's an optimal sub-tree that looks like this; aka, what our algorithm did okay.

- Our algorithm would do this at level t :




We've got a bunch of current sub-trees:



say that x and y are the two smallest.


Inductive step

- Suppose that the inductive hypothesis holds for $t-1$
 - After $t-1$ steps, there is an optimal tree containing all the current sub-trees as “leaves.”




Lemma 2 again says that there's an optimal tree that looks like this

- Our algorithm would do this at level t :




We've got a bunch of current sub-trees:



say that x and y are the two smallest.

Inductive step


- Suppose that the inductive hypothesis holds for $t-1$
 - After $t-1$ steps, there is an optimal tree containing all the current sub-trees as “leaves.”



Lemma 2 again says that there's an optimal tree that looks like this

aka, there is an optimal tree containing all the level- t sub-trees as “leaves”.

- Our algorithm would do this at level t :




This is what we wanted to show for the inductive step.

Inductive outline:

After the t' th step, we've got a bunch of current sub-trees:

- Inductive hypothesis:

- after the t' th step,
 - there is an optimal tree containing the current subtrees as “leaves”




- Base case:

- after the 0'th step,
 - there is an optimal tree containing all the vertices.

Inductive hyp. asserts that our subtrees can be assembled into an optimal tree:


- Inductive step:

- **TO DO**




- Conclusion:

- after the last step,
 - there is an optimal tree containing this whole tree as a subtree.
- aka,
 - after the last step the tree we've constructed is optimal.



What have we learned?

- ASCII isn't an optimal way* to encode English, since the distribution on letters isn't uniform.
- Huffman Coding is an optimal way!
- To come up with an optimal scheme for any language efficiently, we can use a **greedy algorithm**.
- To come up with a **greedy algorithm**:
 - Identify **optimal substructure**
 - Find a way to make choices that **won't rule out an optimal solution**.
 - Create subtrees out of the smallest two current subtrees.



Recap I

- Greedy algorithms!
- Three examples:
 - Activity Selection
 - Scheduling Jobs
 - Huffman Coding



Recap II



- Greedy algorithms!
- Often easy to write down
 - But may be hard to come up with and hard to justify
- The natural greedy algorithm may not always be correct.
- A problem is a good candidate for a greedy algorithm if:
 - it has optimal substructure
 - that optimal substructure is **REALLY NICE**
 - solutions depend on just one other sub-problem.

Next time

- Greedy algorithms for **Minimum Spanning Tree!**

Before next time

- Pre-lecture exercise: thinking about MSTs